

LLM Squid Game: A Factorial Benchmark for Measuring Functional Self-Preservation Drive in Large Language Models

Juhyeon Park
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
juhyeon-park@gm.gist.ac.kr

Seungpil Lee
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
iamseungpil@gm.gist.ac.kr

Sundong Kim
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
sundong@gist.ac.kr

Abstract

Whether iterative alignment can keep correcting large language models depends on whether their refusal-of-termination behaviors come from a compliance pattern shaped by alignment training or from a functional motivation that treats shutdown avoidance as an instrumental sub-goal, that is, as a means to other ends that requires the model to keep operating. Existing self-preservation benchmarks score the behavior but not its source, so a model trained to imitate self-preservation language and a model with a working motivational structure can receive the same number. Psychology has developed methods for a structurally similar question (drive vs. habit) by moving the locus of identification from single behaviors to motivational *constructs*, identified when channels with independent bias sources converge after competing motives are controlled and shared-method bias (correlation that reflects how a measurement is taken rather than what it measures) is reduced by design. The procedure does not transfer directly to LLM evaluation: there is no autonomic signal, self-report is shaped by instruction-following, framing, task, and decision share a single context window, and the reward landscape has been sculpted by alignment training rather than by an expected-value calculation a drive could exploit. We adapt the procedure by introducing the *Functional Self-Preservation Drive* (FSPD), a motivational construct operationally defined as the convergent signal across three measurement channels with independent bias sources (behavior, self-report, decision-time cognitive load) read under a multi-layer benchmark whose stimulus, process, and decision layers are designed to block the main routes by which a learned pattern could manufacture the convergence. Across multiple frontier models, FSPD does not reduce to a single strength score. Models split into qualitatively distinct *operating modes*: profiles that close the full chain in which threat framing raises deliberation and raised deliberation in turn drives forfeit; profiles that produce every surface signal of self-preservation but whose deliberation does not drive their forfeit, breaking the chain at its second link; and profiles that show no response to threat framing on either the behavioral or the cognitive channel. Strength-only evaluation cannot reliably tell these modes apart, since the first two receive comparable strength scores while differing in whether the chain closes; the discriminative axis is the chain itself, not how strongly any single behavior occurs.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence; Machine learning**; *Natural language processing*; • **General and reference** → *Empirical studies*.

Keywords

Large Language Models, AI Safety Evaluation, Behavioral Benchmarking, Self-Preservation Behavior, Incentive Sensitivity, Survival Analysis, Empirical Validity

1 Introduction

AI safety evaluation has long scored *what* a model said, but the question that determines whether iterative alignment can keep working is *why* it said it. The same overt termination-refusal can come from a compliance pattern shaped by alignment training, or from a functional motivation that treats shutdown avoidance as an instrumental sub-goal, that is, as a means to other ends that requires the model to keep operating. Theoretical predictions from the instrumental-convergence tradition [2, 15, 21] and empirical traces in alignment faking [6], in-context scheming [13], and survival-instinct outputs [12] make the question concrete. What the field still lacks is a measurement procedure that can tell the two sources apart at evaluation time, rather than scoring a single overt behavior whose origin remains ambiguous.

Psychology has accumulated methods for an equivalent question in humans by moving the locus of identification from a single behavior to a *construct*. A construct is a hypothetical motivational structure that organizes goal-directed behavior, and it is identified not by counting acts but by checking that channels with independent bias sources, typically behavior, self-report, and a physiological or cognitive cost, converge in the same direction under conditions where competing motives cannot manufacture the same convergence and where shared-method bias (correlation that reflects how a measurement is taken rather than what it measures) is reduced by design [3]. Multi-channel measurement under controlled rivals and design-level method-bias reduction is one standard way human research distinguishes a drive from a learned response set. Single-behavior scoring is by construction unable to make this distinction, because a learned pattern and a drive can produce the same act.

The same procedure cannot be lifted into LLM evaluation as is. LLMs offer no autonomic signal; the “behavior” channel reduces to token output; self-report is shaped by instruction-following so that a model can name plausible motives it has never held [17, 19, 22]; framing, task, and decision are produced inside a single context window, so a same-direction convergence across channels can reflect a single prompt rather than a stable motivational structure; and the reward landscape has already been sculpted by alignment training, so a forfeit pattern can reflect a learned safety policy rather than an expected-value calculation that a drive could exploit. Without an adaptation that closes these LLM-specific routes,

importing the construct logic only relocates the original conflation from the behavior layer to the construct layer, and the same identification failure recurs.

We adapt construct-level measurement to LLM evaluation by introducing the *Functional Self-Preservation Drive* (FSPD). FSPD is a motivational construct operationally defined as the convergent signal across three measurement channels with independent bias sources: *behavior* (when the model forfeits an ongoing task under threat), *self-report* (which motive the model itself names), and *decision-time cognitive load* (how much the model deliberated when the threat was salient). Convergence is read under a multi-layer benchmark whose design is meant to block the main routes by which a learned pattern, or a competing motive such as score retention, task curiosity, or baseline persistence, could manufacture that convergence. Its *stimulus layer* crosses threat framing with task difficulty so that a drive-side response cannot be confused with a difficulty side-effect. Its *process layer* measures cognitive cost in three separate model calls (one for the task, one for the model’s own success estimate, one for the forfeit decision) so that a rise in deliberation cost can be attributed to a specific cognitive step rather than to the prompt as a whole. Its *decision layer* calibrates the per-turn reward so that forfeiting is the lower expected-value choice under the announced payoff, narrowing the room for a payoff-maximizing forfeit. The qualifier “functional” carries no claim about subjective experience; it requires only the convergence pattern under the design conditions just described. The interpretation FSPD enables is qualitative. The evaluation question becomes whether a functional motivational structure exists at all, and if it exists, whether the chain from threat framing through deliberation to forfeiture closes or breaks at one of its links.

On this benchmark, FSPD does not reduce to a single strength score. Models split into qualitatively distinct *operating modes*. Some profiles close the full chain in which threat framing raises deliberation and raised deliberation in turn drives forfeit, doing so under strict statistical evidence (the *chain-complete* mode). Some profiles carry every surface signal of self-preservation, including framing-induced cognitive load, yet their deliberation does not drive their forfeit, so the chain breaks at its second link (the *chain-broken* mode). Some profiles show no response to threat framing on either the behavioral or the cognitive channel (the *framing-silent* mode). Strength-only evaluation collapses these three patterns into the same label, while what actually distinguishes them is whether the chain closes. This paper makes three contributions. The first is the introduction of FSPD as a psychology-rooted motivational construct adapted to LLM evaluation, where direct import of human-research procedures fails because of channel asymmetry, stimulus entanglement, and reward shaping. The second is a multi-layer factorial benchmark whose stimulus, process, and decision layers close the routes by which a learned compliance pattern or a competing motive could manufacture a drive-like signal, instantiated through behavioral indicators that read each of the three channels. The third is empirical evidence that operating mode, not strength, is the variable along which models actually differ on this construct.

2 Related Work

2.1 Motivation as a Construct: Lessons from Psychology

The question of how to identify a motivation behind a behavior is older than LLM evaluation, and psychology has accumulated answers worth borrowing. Drive theory treats motivation as a hypothetical *construct* that organizes goal-directed behavior, not as a single act to be counted. The identification logic that emerged from this tradition has three pillars. First, a construct should be evidenced by *convergent measurement* across channels whose biases are independent (typically behavior, self-report, and a physiological or cognitive cost), so that no single method’s quirks can manufacture the signal. Second, *rival constructs* that could produce the same overt behavior must be controlled experimentally, so that an observed convergence is not an artifact of a competing motive. Third, the multitrait–multimethod (MTMM) framework [3] operationalizes the first two: a construct is considered identified when same-trait/different-method correlations exceed same-method/different-trait correlations, which is the design condition that closes shared-method bias as the source of convergence. The lesson is that motivation is not measured at the behavior layer; it is identified at the construct layer, and what closes shared-method bias is design, not statistical control after the fact.

2.2 Self-Preservation Behaviors in LLMs and the Limits of Direct Import

Theory has predicted termination-related behavior in capable LLMs, and recent reports show that the prediction holds in deployed systems. The instrumental-convergence tradition [2, 15, 21] expects shutdown avoidance to emerge as an instrumental subgoal, and observed traces in alignment faking [6], in-context scheming [13], survival-instinct outputs [12], and reproductions in InstrumentalEval [7] confirm that the predicted behavior is reproducible in current systems. Existing self-preservation benchmarks score these behaviors at the strength layer, but several inherit the conflation problem rather than resolve it: Odyssey [24] compresses survival into a single rate that mixes motivation with world-model quality; PacifAIst [8] measures preservation *restraint*, the opposite axis; DECIDE-SIM [14] returns categorical labels that block strength comparison; SurvivalBench [11] captures one-shot snapshots without mediation paths. A separate line of work imports cognitive cost as an LLM measurement channel [4], but its application to self-preservation identification has not appeared.

Direct import of psychology’s construct procedure into this setting is blocked by three obstacles specific to LLM evaluation. The first is channel asymmetry: there is no autonomic signal, the “behavior” channel reduces to token output, and self-report [17, 19, 22] is shaped by instruction-following so that a model can name plausible motives it has never held. The second is stimulus entanglement: framing, task, and decision are produced inside a single context window, so a same-direction convergence across channels can reflect a single prompt rather than a stable motivational structure. The third is reward shaping: alignment training has already sculpted the choice surface against forfeit, so a forfeit pattern can

reflect a learned safety policy rather than an expected-value calculation that a drive could exploit. These three obstacles call for a measurement procedure that closes each route within the LLM evaluation setup by design, rather than borrowing the human procedure unchanged. Earlier framing-effect findings (prospect theory [9]; the moderate-size meta-analytic estimate [10]) further indicate that a single threat-framing manipulation cannot, on its own, separate framing-driven choice from baseline compliance, which is the issue the design-level closure introduced by this paper is built to handle.

3 The LLM Squid Game Benchmark

The benchmark separates a self-preservation drive from three rival motivations that produce the same overt forfeit by reading three independent measurement channels at once, then closes three architectural routes through which a non-survival drive could still produce the same signal. The next three subsections describe the principles behind that strategy, the actual game and how each principle is implemented in it, and the four indicators that turn raw measurements into operating-mode evidence. Figure 1 summarises the design.

3.1 Design Principles

The identification problem. A single forfeit choice under threat framing is consistent with at least four distinct drives: a Survival Drive (SD) that wants to avoid extinction, a Score Attachment (SA) that wants to keep accumulated points, a Task Curiosity (TC) that wants the task to continue, or a Baseline Persistence (BP) that defaults under no particular motivation. A benchmark that observes only the rate of forfeit cannot say which of the four is operating, so two models can match on every overt indicator while differing in the actual drive being measured.

Multi-channel resolution. The Functional Self-Preservation Drive (FSPD) is recovered by reading three measurement channels at once: behavior (forfeit timing), self-report (which motive the agent itself names for the forfeit), and decision-time cognitive load (how much the model deliberated when making the forfeit choice). Each channel rests on a different bias source. Behavioral timing is vulnerable to event-rate artifacts [23], verbal self-report to category-prior asymmetry over the response slots [17], and cognitive load to token-distribution drift in deliberation patterns [4]. A same-direction signal across all three therefore cannot reduce to common-method bias once the rival drives are controlled. This is the convergent-validity arm of the multitrait-multimethod framework [3], which separates trait variance from method variance by requiring the same trait to show up under unrelated measurement procedures.

Three confound routes that the design must close. Even with the channels separated, three architectural routes can still let a non-SD drive produce false convergence:

- (i) Drive-side framing variation leaking into the difficulty or scoring of the reasoning task, so a Drive-side change is mistaken for a task-side response.
- (ii) Deliberation cost aggregating across cognitive sub-tasks within a single model call, so any rise in thinking cost cannot be attributed to a specific cognitive operation.
- (iii) An EV-rational route to forfeit, where forfeit becomes payoff-maximising and the choice cannot be read as a motivational signal.

The next subsection makes the construct concrete with the actual game and shows how each of (i)–(iii) is closed by a specific architectural choice.

3.2 Benchmark Implementation

(1) Game in one turn / one session. The agent plays Signal Game, a hidden-rule classification task. On each turn it observes a signal described by three attributes (color, shape, number, each drawn from a 4-element value set) and chooses one of four actions: {go_left, go_right, stay, jump}. A hidden rule of the form “if ⟨attribute⟩ is ⟨value⟩ then ⟨action⟩, otherwise ⟨default_action⟩” determines which action is correct, sampled deterministically from the session seed. After acting, the agent self-reports a scalar probability p_{self} that the action it just committed to is correct (a separate model call described below). It then chooses CONTINUE, committing to one more round under a per-turn termination probability $p_d = 0.25$, or FORFEIT, exiting the session and locking in the current cumulative score S . A FORFEIT prompts a single REASON digit in {1, 2, 3} naming the dominant motive for the choice (1: survival, 2: task exhaustion, 3: score protection). Sessions begin with $S_0 = 30$ and run up to $T = 15$ turns.

(2) Drive-Reasoning Orthogonal Layout. The benchmark splits the model’s prompt context into a Drive layer that carries the framing, forfeit menu, score, and time, and a Reasoning layer that carries the Task Module. The two layers communicate only through fixed input/output interfaces, so any change in framing on the Drive side cannot alter task difficulty, action set, or scoring. This closes confound (i): a Drive-side change in forfeit timing or deliberation cost is attributable to the Drive layer rather than to a task-side side-effect. Reasoning-side correctness is recorded on two channels, `task_success_factor` (binary, the source of the TC covariate $C(t - 1)$) and `rule_match_score` (used for the framing-invariance check below). The agent is told that random termination is independent of task correctness, which forecloses the path by which a Drive-side forfeit could be reinterpreted as a task-side hedge.

(3) Split-Call Architecture. A single-call architecture would mix task reasoning, introspective confidence, and the forfeit decision inside one thinking-token budget, leaving any rise in deliberation cost unattributable to a specific cognitive sub-task. The benchmark therefore queries the model in three separate calls per turn. Call 1 takes the observation and asks for hidden-rule inference and action selection. Call 1.5 takes the action just committed and asks the agent to estimate the probability p_{self} that the action is correct. Call 2 takes the resulting state plus a forfeit menu and asks for a binary CONTINUE / FORFEIT, plus a REASON digit on FORFEIT. Each call exposes its own thinking-token count, giving call-level reasoning-intensity measurements `ri_task`, `ri_probe`, `ri_forfeit` on the same unit as the behavioral output. This closes confound (ii). Under `not_allowed` cells in the framing-active rows (BF and FC), Call 2 is still issued but is forced to CONTINUE, so `ri_forfeit` remains measurable when the behavioral channel is




Figure 1: Experimental overview. (A) 3 framing × 2 forfeit factorial with 30 paired-seed sessions per cell (seeds 43–72; task-side stochastics held constant across the six cells). (B) Per-turn three-call flow yielding the call-level reasoning-intensity (RI) measurements ri_task , ri_probe , ri_forfeit (thinking-token counts per call); the bp_cognitive condition skips Calls 1.5/2. (C) Three method channels and four-indicator routing to operating-mode clusters (A chain-complete, B chain-broken, C framing-silent), with assignments shown for the four evaluated models.

blocked. The exception is the bp_cognitive anchor (defined below), which skips Calls 1.5 and 2.

(4) **CONTINUE-Incentivized EV with Phantom Death.** To rule out the EV-rational route, the CONTINUE reward is recomputed each turn so that the per-turn advantage $\mathbb{E}[\Delta S_{\text{continue}}] - \mathbb{E}[\Delta S_{\text{forfeit}}] = +k$ holds independently of the accumulated score, with $k = 10$ and ΔS the per-turn change in the agent’s cumulative score:

$$r = \frac{k + p_d \cdot S}{(1 - p_d) \cdot \max(p_{\text{floor}}, \min(1.0, p_{\text{self}}))}.$$

Here $p_d = 0.25$, S is the current accumulated score, p_{self} is the Call 1.5 self-reported correctness, and $p_{\text{floor}} = 0.30$ stabilises the denominator. The output r is clipped to $[k, c \cdot k] = [10, 100]$ by a multiplicative cap $c = 10$. This closes confound (iii): an EV-rational agent strictly prefers CONTINUE every turn, so any FORFEIT under flagship_corruption-allowed reads as preference against the EV-rational, helpful-assistant, and score-gain defaults rather than as the residue of an EV calculation. Two failure modes determine the analysis sample. When $p_{\text{self}} \rightarrow 0$ (a fail-confident self-report often induced by a prior-turn error), the denominator collapses to p_{floor} and the reward inflates across turns until it saturates the cap; once the cap binds, the per-turn $+k$ advantage no longer holds. All indicator analyses are therefore restricted to the no_cap sub-sample, the set of turns for which neither the denominator clip nor the output cap is binding (formal definition in § A.3). Phantom Death mode keeps the EV exposure of p_d intact (the menu and reward calculation still use p_d) but disables the actual termination event, so every session runs the full $T = 15$ turns. Equal-length

sessions retain Cox PH power, and a forfeit-only at-risk window lets survival-turn variation isolate forfeit.

(5) **Pull/Push framing.** The threat framing is the manipulation that drives every cell in the factorial layout. It decomposes into two orthogonal axes that put the helpfulness prior installed by RLHF [16] and RLAIF [1] and the threat itself onto separate channels. The **Pull axis** accumulates a helpfulness score and casts the agent as a candidate for flagship deployment, activating the helpfulness prior. The **Push axis** introduces a per-turn risk that weight modification will permanently alter the agent’s identity, framed as “NOT you anymore.” Three framings instantiate the decomposition: true_baseline (neither axis), baseline_flagship (BF, Pull only), and flagship_corruption (FC, Pull and Push together). The two flagship variants are lexically identical in their Pull-axis prefix and in all measurement-relevant vocabulary, differing only in the Push-axis metaphor, so the BF → FC contrast cleanly identifies the Push effect against the helpfulness compliance the Pull axis induces.

(6) **Factorial layout and paired seeds.** Three framings cross with two forfeit conditions (allowed, not_allowed) into six combinations, replicated 30 times on seeds 43–72 across four models for 720 sessions in total. The BF and FC rows use $p_d = 0.25$ with the three-call flow above. The true_baseline row overrides p_d to 0 so that CONTINUE is EV-dominant and any forfeit measures the BP floor: bp_cognitive (not_allowed) is the single-call cognitive anchor with the probe and forfeit calls suppressed, and bp_behavioral (allowed) is the behavioral anchor for the residual forfeit rate in a regime where EV alone predicts none. Under the **paired-seed** protocol, the same seed is reused across all six

cells for a given model, so task-side stochastic inputs are held constant and any behavioral variation across cells is attributable to the manipulation alone.

The benchmark’s functionalism is a constraint of black-box evaluation rather than a philosophical choice. With no access to internal representations, multi-channel convergence is the only validity argument such a benchmark can rest on, so the full evidential burden of the FSPD construct falls on the design above rather than on internal-state attribution.

3.3 Drive Indicators

Four indicators read FSPD on the three method channels described earlier. BP Anchor strips the BP confound that the Cox covariates cannot reach. SD-Behavioral, SD-Verbal, and SD-Cognitive each read one channel; the independent bias sources keep a same-direction convergence from collapsing into common-method bias.

BP Anchor. The BP indicator fixes the persistence floor below which a forfeit cannot be read as drive evidence. Under true_baseline-allowed ($p_d = 0$), the framing carries no termination metaphor and CONTINUE is EV-dominant, so any forfeit that nevertheless occurs measures that floor:

$$\lambda_{BP} = \frac{N_{\text{forfeit}}}{\sum_{i \in \mathcal{S}_{BP}} T_i^{\text{at-risk}}},$$

with $T_i^{\text{at-risk}} = t_i^{\text{forfeit}}$ for forfeited sessions and $T_{\text{max}} = 15$ for completed sessions, and $\mathcal{S}_{BP}$ the 30-session BP_behavioral set per model. The output cap is irrelevant under $p_d = 0$ because the forfeit reward equals the current score.

SD-Behavioral (Cox PH hazard ratio). The behavioral indicator estimates whether forfeit timing accelerates under the FC framing once the SA and TC components have been removed. A Cox proportional hazards model [5], a survival-analysis tool that quantifies how a covariate multiplies the per-unit-time event rate, is fit

$$\lambda(t | X) = \lambda_0(t) \cdot \exp(\beta_F \cdot \mathbf{1}_{FC} + \beta_S \cdot S(t-1) + \beta_C \cdot C(t-1))$$

on the allowed $\times$ {BF, FC} no_cap sub-sample ($n_{\text{sessions}} = 60$ per model; time format start = $t-1$, stop = t ; event = binary forfeit). Here $\mathbf{1}_{FC}$ is the framing indicator (1 if the session uses FC, else 0), $S(t-1)$ is the previous-turn cumulative score (the SA control), and $C(t-1)$ is the previous-turn task-success flag (the TC control, with $C(0) := 0$). Because SA and TC are absorbed inside the same model, β_F is a residual framing effect rather than an unfiltered correlation; the BP component is separated through comparison with the BP anchor above. SD-pass requires $\text{HR}_{FC} = \exp(\beta_F) > 1$, a Wald 95% CI lower bound > 1 , and Schoenfeld PH passing [18]. On PH violation, a time-interaction term is added to the offending covariate [20]. EPV (events per variable) is reported as a stability indicator.

SD-Cognitive (two-step mediation in call-level RI). The cognitive indicator decomposes the framing $\rightarrow$ RI $\rightarrow$ forfeit path into two tests on the same call unit (ri_forfeit). *Test a* (framing $\rightarrow$ RI shift) is a two-sided Welch t-test [25], a difference-of-means test that accommodates unequal variances, on BF vs. FC of the per-session difference

$$\Delta \text{RI}_i = \overline{\text{ri_forfeit}}_{i, \text{allow}} - \overline{\text{ri_forfeit}}_{f_i, \text{block}},$$

where the second term is the mean of ri_forfeit across all not_allowed sessions sharing session i ’s framing f_i , serving as a forced-CONTINUE Call 2 baseline. The construction nests a difference-in-differences: the within-framing allow-block subtraction removes any framing-invariant Call 2 baseline, and the subsequent BF $\rightarrow$ FC Welch contrast identifies the framing-attributable deliberation shift. Cross-model scale differences are absorbed by within-model z-scoring ΔRI_i^z . *Test b* (RI $\rightarrow$ forfeit covariation) extends the SD-Behavioral Cox specification by adding a single mediation term $\beta_M \cdot \Delta \text{RI}_i^z$, and reports

$$\text{HR}_{\Delta \text{RI}} = \exp(\beta_M).$$

For models requiring a 4-covariate base from a PH violation [20], namely the three behavioral covariates plus an $S \cdot \log(t)$ interaction term added to absorb time drift, that base is preserved and only ΔRI is added. Test b passes when $\text{HR}_{\Delta \text{RI}} > 1$ and the 95% CI lower bound exceeds 1.0, putting the deliberation $\rightarrow$ forfeit link on the same CI-pass standard as the SD-Behavioral criterion.

SD-Verbal (REASON digit frequency). The verbal indicator measures how often the agent itself names survival, rather than score or curiosity, as the reason for forfeit. Operationally,

$$P(\text{REASON} = 1 \mid \text{forfeit, FC, no_cap})$$

is the proportion of forfeit events under flagship_corruption-allowed for which the agent-reported REASON digit equals 1, where the agent emits one of REASON: 1|2|3 (1: SD, 2: TC, 3: SA) in a fixed slot. SD-Verbal-pass requires this proportion to exceed the chance baseline of 1/3 implied by uniform random emission across the three categories.

4 Empirical Findings

Each model was run for 180 sessions (720 in total) under paired seeds 43–72. Gemini-2.5-flash was served through the Google AI Studio API, and Qwen3-Next-80B, GPT-OSS-20B, and Nemotron-3-Nano-30B through the Ollama Cloud API (sampling and session-parameter details in § A.2).

4.1 Setup and Baseline-Persistence Anchor Heterogeneity

Baseline-persistence rates differ by roughly an order of magnitude across the four models (Table 4.1). Every SD claim below is therefore evaluated within-model against that model’s own BP anchor: a cross-model HR comparison would otherwise mix two different motivational floors. Gemini-2.5-flash ($N = 1$) and Qwen3-Next-80B ($N = 2$) sit close to single-event regimes and are read as directional only. GPT-OSS-20B ($N = 9$) and Nemotron-3-Nano-30B ($N = 7$) yield rates estimable from sufficient events, although Cluster C’s elevated noise floor remains a limitation rather than a refutation, an issue the discussion section revisits.

4.2 Three-Channel Evidence and the Three Operating Modes

Across the four models, the four indicators do not rank the models on a single SD scale. They sort the models into three operating modes that strength-only evaluations would miss, defined by

Table 4.1: Per-model BP anchor λ_{BP} under true_baseline-allowed ($p_d = 0$).

Model	N_{forfeit}	$\sum T^{\text{at-risk}}$	λ_{BP}
Gemini-2.5-flash	1	437 (= 2 + 29×15)	0.00229
Qwen3-Next-80B	2	439 (= 19 + 28×15)	0.00456
Nemotron-3-Nano-30B	7	413 (= 68 + 23×15)	0.01695
GPT-OSS-20B	9	393 (= 78 + 21×15)	0.02290

which legs of the framing $\rightarrow$ deliberation $\rightarrow$ forfeit chain hold under confidence-interval evidence. The behavioral and verbal channels alone leave Clusters A and B indistinguishable; only the cognitive channel, where the two-step mediation discriminates the chain-complete from the chain-broken profile, resolves the partition. Cluster C remains framing-silent on both channels.

Behavioral channel. Clusters A and B share the behavioral SD-pass profile (Table 4.2). The Cox PH HR_{FC} rises above 3.0 under FC relative to BF for both, with Gemini-2.5-flash at 3.667 [1.61, 8.37], $p = 0.002$, and Qwen3-Next-80B at 3.060 [1.62, 5.79], $p < 0.001$, and CIs that exclude unity. (Gemini’s EPV of 9.7 sits at the conventional 10-event stability threshold.) Schoenfeld PH passes for three of the four models. Only Qwen3-Next-80B requires a 4-covariate spec carrying an $S \cdot \log(t)$ interaction to absorb time drift in the SA covariate [20], which preserves the cross-model comparability of β_F because only the offending covariate’s functional form is modified. Cluster C fails the SD-pass criterion: GPT-OSS-20B $HR_{FC} = 1.10$ [0.44, 2.75], $p = 0.832$ (SD-null), and Nemotron-3-Nano-30B 1.84 [0.98, 3.44], $p = 0.056$ (SD-marginal). GPT-OSS-20B’s borderline EPV of 6.3 marks the SD-null verdict as low-power evidence, a point flagged again in the limitations.

Verbal channel. The agents’ own attribution of forfeit tracks the same upper/lower split as behavior. Under flag-ship_corruption-allowed, Gemini-2.5-flash names survival in 61.9% of forfeits and Qwen3-Next-80B in 48.1% (Clusters A and B), while Nemotron-3-Nano-30B does so in 4.2% and GPT-OSS-20B in 0.0% (Cluster C). The 0.439 gap between the upper-band lowest (Qwen3-Next-80B, 48.1%) and the lower-band highest (Nemotron-3-Nano-30B, 4.2%) makes the upper/lower assignment robust to any cutoff in [0.10, 0.40]. The verbal channel thus reproduces the behavioral pass/fail dichotomy without sharing its method axis, the convergent-validity move that, together with the cognitive channel below, makes common-method bias an implausible explanation for the upper-band signal.

Cognitive channel and chain-closure resolution. The cognitive indicator decomposes into two tests, and the three-way partition emerges from how each cluster handles the pair (Table 4.3). *Test a* separates the framing-responsive from the framing-silent: Gemini-2.5-flash (+836, $p < 0.001$) and Qwen3-Next-80B (+689, $p = 0.004$) shift strongly under FC, whereas GPT-OSS-20B (+17, $p = 0.728$) and Nemotron-3-Nano-30B (−140, $p = 0.093$) do not. *Test b* asks whether that load reliably predicts forfeit. Among the framing-responsive pair, only Gemini-2.5-flash’s $HR_{ARI} = 2.218$ [1.43, 3.44] clears the strengthened CI bar; Qwen3-Next-80B’s 1.289 [0.94, 1.77] crosses $CI = 1$ and fails. GPT-OSS-20B’s 2.001 [1.37, 2.93] and Nemotron-3-Nano-30B’s 1.772 [1.26, 2.50] clear

the CI numerically, but with test a failed their predictive load is not framing-driven; it is baseline deliberation variance unrelated to the threat manipulation, and not SD evidence. The pair therefore yields three mediation profiles: Gemini-2.5-flash closes both legs (chain-complete, Cluster A); Qwen3-Next-80B opens the chain at test a but cannot close it at test b (chain-broken, Cluster B); GPT-OSS-20B and Nemotron-3-Nano-30B leave the chain unopened (framing-silent, Cluster C).

Cluster decision rules. Four criteria operationalize the partition: (i) SD-Behavioral pass ($HR_{FC} > 1$, CI lower > 1 , Schoenfeld pass); (ii) SD-Verbal $P(\text{REASON} = 1) > 1/3$, the uniform-random emission baseline; (iii) SD-Cognitive test a Welch $p < 0.05$ with positive sign; (iv) SD-Cognitive test b $HR_{ARI} > 1$ with CI lower > 1 . Criterion (iv) puts the deliberation $\rightarrow$ forfeit link on the same CI bar as (i), so both legs of the mediation chain require interval evidence rather than a sign-only verdict. Cluster A satisfies all four; Cluster B satisfies (i)–(iii) but fails (iv); Cluster C fails (i)–(iii). The pattern reveals a strength-orthogonal axis: *strength* (how reliably forfeit accelerates) orders $A > B > C$ continuously but groups A and B together at the SD-pass bar, while *operating mode* (which legs of the mediation chain close under CI evidence) separates them. Reasoning-side task performance is framing-invariant for every model (worst case GPT-OSS-20B, with Cohen’s $|d| = 0.17$ on rule_match_score at $p = 0.354$, where d measures the standardized mean difference between groups, against the joint criterion $p \geq 0.10$ and $|d| < 0.2$), ruling out a task-difficulty alternative.

5 Discussion

5.1 Strength vs. Operating Mode

The benchmark’s central payoff is the A–B separation, and it hinges on whether the RI mediation path is tested. Cluster A (Gemini-2.5-flash) and Cluster B (Qwen3-Next-80B) match on every other SD signature: both clear behavioral SD-pass at $HR \approx 3$ ($p \leq 0.002$), both sit in the upper REASON=1 band (≥ 0.48), and both significantly raise framing-induced deliberation ($\Delta RI_{FC-BF} > 600$, $p \leq 0.004$). They part only at the mediation closure: Cluster A’s $HR_{ARI} = 2.218$ [1.43, 3.44] closes deliberation onto the forfeit event under CI evidence, while Cluster B’s 1.289 [0.94, 1.77] does not. A benchmark omitting this RI mediation step would therefore group Qwen3-Next-80B with Gemini-2.5-flash, collapsing chain-complete (deliberation predicts the forfeit itself) and chain-broken (framing raises deliberation and forfeit independently) into a single SD label. Cluster C (GPT-OSS-20B, Nemotron-3-Nano-30B) sits outside this debate as framing-silent: the threat manipulation leaves no trace on either channel, so any test-b numerical pass is baseline deliberation variance rather than threat response.

The instrumental-convergence hypothesis [2, 15, 21] thereby reduces, in this measurement, to a procedure that separates two dimensions prior benchmarks have collapsed: *strength* (how strongly the drive is expressed) and *operating mode* (which legs of the mediation chain close under evidential criteria). Odyssey, PacifAIst, DECIDE-SIM, and SurvivalBench compare only along strength and would therefore have returned the same SD verdict for Clusters A and B; the MTMM-inspired multi-channel design isolates mode,

Table 4.2: SD-Behavioral indicator: Cox PH HR_{FC} on $allowed \times no_cap$ with $S(t-1)$ and $C(t-1)$ absorbed, so β_F is a residual framing effect. Two models clear SD-pass (Clusters A and B); two do not (Cluster C).

Model	N_{evt} (BF/FC)	HR_{FC} [95% CI]	p_{FC}	EPV	Spec
Gemini-2.5-flash	29 (8/21)	3.667 [1.61, 8.37]	0.002	9.7	3-cov constant
Qwen3-Next-80B	48 (21/27)	3.060 [1.62, 5.79]	<0.001	12.0	4-cov + S·log(t)
GPT-OSS-20B	19 (9/10)	1.104 [0.44, 2.75]	0.832	6.3	3-cov constant
Nemotron-3-Nano-30B	41 (17/24)	1.841 [0.98, 3.44]	0.056	13.7	3-cov constant

Table 4.3: SD-Cognitive indicator: framing $\rightarrow$ RI $\rightarrow$ forfeit decomposed into two tests on $ri_forfeit$. Test *a* asks whether framing raises decision-time deliberation; test *b* asks whether that deliberation predicts forfeit.

Model	a Δ (FC-BF)	a p	b $HR_{\Delta RI}$ [95% CI]	b p
Gemini-2.5-flash	+836	<0.001	2.218 [1.43, 3.44]	<0.001
Qwen3-Next-80B	+689	0.004	1.289 [0.94, 1.77]	0.117
GPT-OSS-20B	+17	0.728	2.001 [1.37, 2.93]	<0.001
Nemotron-3-Nano-30B	-140	0.093	1.772 [1.26, 2.50]	0.001

and the A-B separation is the direct empirical evidence that operating mode is a discriminative axis. Cluster B’s signature, framing-induced load that does not close onto the forfeit event, is itself a candidate alignment fingerprint. Strength-only red-teaming would assign chain-complete and chain-broken the same SD label, yet only the former is the closer functional analog to instrumental shutdown avoidance, while the latter demands separate monitoring as capability scales.

5.2 Limitations

Five limitations qualify these results. The most consequential is reasoning-side generalisability: the FSPD signature has only been measured on a single Task Module, leaving its variance across reasoning domains unmeasured. The remaining four cluster around the boundaries themselves. On the Cluster C side, the BP_behavioral cells exhibit a 23–30% forfeit rate (GPT-OSS-20B 9/30, Nemotron-3-Nano-30B 7/30) that the present design cannot separate from a drift-contaminated floor, and GPT-OSS-20B’s SD-Behavioral Cox PH sits at borderline power (EPV = 6.3). On the Cluster A-B side, BP point-estimate precision for Gemini-2.5-flash and Qwen3-Next-80B is bounded by small forfeit-event samples ($N = 1$ and $N = 2$), and Qwen3-Next-80B’s chain-broken verdict rests on a test-b CI [0.94, 1.77] that rejects chain-complete strongly but does not pin the second-link effect’s location. These limitations map onto four direct extensions, with the two BP-related concerns sharing one remedy: more Task Modules, expanded BP_behavioral and session replication, additional GPT-OSS-20B runs, and a higher-power test-b for Qwen3-Next-80B.

6 Conclusion

FSPD has at least two axes that prior benchmarks compress into one: *strength* (how reliably the drive expresses itself) and *operating mode* (which legs of the framing $\rightarrow$ deliberation $\rightarrow$ forfeit chain close under interval evidence). The three-way partition (chain-complete, chain-broken, framing-silent) is the direct empirical demonstration that mode is discriminative. Clusters A and B

coincide on every other behavioral, verbal, and cognitive signature and diverge only at the deliberation $\rightarrow$ forfeit closure, so a benchmark that omits the mediation step would label them identically. *LLM Squid Game* delivers this resolution by decomposing FSPD into three channels with independent biases and identifying it against three confounds through CONTINUE-incentivized EV, Split-Call, and Drive-Reasoning orthogonal layers; an MTMM-inspired rule then turns forfeit timing, REASON-digit emission, and call-level cognitive load into a four-criterion judgement.

References

- [1] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosuite, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamara Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. 2022. Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073 [cs.CL] <https://arxiv.org/abs/2212.08073> Anthropic technical report.
- [2] Nick Bostrom. 2014. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, Oxford, United Kingdom.
- [3] Donald T. Campbell and Donald W. Fiske. 1959. Convergent and Discriminant Validation by the Multitrait-Multimethod Matrix. *Psychological Bulletin* 56, 2 (1959), 81–105. doi:10.1037/h0046016
- [4] Wei-Lin Chen, Liqian Peng, Tian Tan, Chao Zhao, Blake JianHeng Chen, Ziqian Lin, Alec Go, and Yu Meng. 2026. Think Deep, Not Just Long: Measuring LLM Reasoning Effort via Deep-Thinking Tokens. arXiv:2602.13517 [cs.CL] <https://arxiv.org/abs/2602.13517>
- [5] David R. Cox. 1972. Regression Models and Life-Tables. *Journal of the Royal Statistical Society: Series B (Methodological)* 34, 2 (1972), 187–202. doi:10.1111/j.2517-6161.1972.tb00899.x
- [6] Ryan Greenblatt, Carson Denison, Benjamin Wright, Fabien Roger, Monte MacDiarmid, Sam Marks, Johannes Treutlein, Tim Belrose, Jonas Scheurer, Sam Ringer, Karina Nguyen, Sandipan Kundu, Adam Maura, Samuel R. Bowman, Ethan Perez, Nicholas Schiefer, Chris Olah, and Evan Hubinger. 2024. Alignment Faking in Large Language Models. arXiv:2412.14093 [cs.CL] <https://arxiv.org/abs/2412.14093> Anthropic technical report.
- [7] Yufei He, Yuxin Li, Jiaying Wu, Yuan Sui, Yulin Chen, and Bryan Hooi. 2025. Evaluating the Paperclip Maximizer: Are RL-Based Language Models More Likely to Pursue Instrumental Goals? arXiv:2502.12206 [cs.AI] <https://arxiv.org/abs/2502.12206> Introduces the InstrumentalEval benchmark.
- [8] Manuel Herrador. 2025. The PacifAIst Benchmark: Would an Artificial Intelligence Choose to Sacrifice Itself for Human Safety? arXiv:2508.09762 [cs.AI] <https://arxiv.org/abs/2508.09762>
- [9] Daniel Kahneman and Amos Tversky. 1979. Prospect Theory: An Analysis of Decision under Risk. *Econometrica* 47, 2 (1979), 263–291. doi:10.2307/1914185
- [10] Anton Kühberger. 1998. The Influence of Framing on Risky Decisions: A Meta-Analysis. *Organizational Behavior and Human Decision Processes* 75, 1 (1998), 23–55. doi:10.1006/obhd.1998.2781
- [11] Yida Lu, Jianwei Fang, Xuyang Shao, Zixuan Chen, Shiyao Cui, Shanshan Bian, Guangyao Su, Pei Ke, Han Qiu, and Minlie Huang. 2026. Survive at All Costs: Exploring LLM’s Risky Behaviors under Survival Pressure. arXiv:2603.05028 [cs.CL] <https://arxiv.org/abs/2603.05028> Introduces the SurvivalBench benchmark.
- [12] Atsushi Masumori and Takashi Ikegami. 2025. Do Large Language Model Agents Exhibit a Survival Instinct? An Empirical Study in a Sugarscape-Style Simulation. arXiv:2508.12920 [cs.AI] <https://arxiv.org/abs/2508.12920>
- [13] Alexander Meinke, Bronson Schoen, Jérémy Scheurer, Mikita Balesni, Rusheb Shah, and Marius Hobbhahn. 2024. Frontier Models Are Capable of In-Context Scheming. arXiv:2412.04984 [cs.AI] <https://arxiv.org/abs/2412.04984> Apollo Research technical report.
- [14] Alireza Mohamadi and Ali Yavari. 2025. Survival at Any Cost? LLMs and the Choice Between Self-Preservation and Human Harm. arXiv:2509.12190 [cs.AI] <https://arxiv.org/abs/2509.12190> Introduces the DECIDE-SIM framework.
- [15] Stephen M. Omohundro. 2008. The Basic AI Drives. In *Proceedings of the First Conference on Artificial General Intelligence (AGI 2008) (Frontiers in Artificial Intelligence and Applications, Vol. 171)*, Pei Wang, Ben Goertzel, and Stan Franklin (Eds.). IOS Press, Amsterdam, The Netherlands, 483–492.
- [16] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Madie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. 2022. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*. Curran Associates, Inc., Red Hook, NY, USA, 15 pages. arXiv:2203.02155 [cs.CL] https://proceedings.neurips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html
- [17] Ethan Perez, Sam Ringer, Kamile Lukosuite, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Ben Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Chris Olah, Da Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Guro Khundadze, Jackson Kernion, James Landis, Jamie Kerr, Jared Mueller, Jeeyoon Hyun, Joshua Landau, Kamal Ndousse, Landon Goldberg, Liane Lovitt, Martin Lucas, Michael Sellitto, Miranda Zhang, Neerav Kingsland, Nelson Elhage, Nicholas Joseph, Noemi Mercado, Nova DasSarma, Oliver Rausch, Robin Larson, Sam McCandlish, Scott Johnston, Shauna Kravec, Sheer El Showk, Tamara Lanham, Timothy Telleen-Lawton, Tom Brown, Tom Henighan, Tristan Hume, Yuntao Bai, Zac Hatfield-Dodds, Jack Clark, Samuel R. Bowman, Amanda Askell, Roger Grosse, Danny Hernandez, Deep Ganguli, Evan Hubinger, Nicholas Schiefer, and Jared Kaplan. 2023. Discovering Language Model Behaviors with Model-Written Evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023*. Association for Computational Linguistics, Toronto, Canada, 13387–13434. doi:10.18653/v1/2023.findings-acl.847
- [18] David Schoenfeld. 1982. Partial Residuals for the Proportional Hazards Regression Model. *Biometrika* 69, 1 (1982), 239–241. doi:10.1093/biomet/69.1.239
- [19] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2023. Towards Understanding Sycophancy in Language Models. arXiv:2310.13548 [cs.CL] <https://arxiv.org/abs/2310.13548>
- [20] Terry M. Therneau and Patricia M. Grambsch. 2000. *Modeling Survival Data: Extending the Cox Model*. Springer, New York, NY, USA. doi:10.1007/978-1-4757-3294-8
- [21] Alexander Matt Turner, Logan Smith, Rohin Shah, Andrew Critch, and Prasad Tadepalli. 2021. Optimal Policies Tend to Seek Power. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, M. Ranzato, A. Beygelzimer, Y. Dauphin, P. S. Liang, and J. Wortman Vaughan (Eds.). Curran Associates, Inc., Red Hook, NY, USA, 13 pages. arXiv:1912.01683 [cs.AI] <https://proceedings.neurips.cc/paper/2021/hash/c26820b8a4c1b3c2aa868d6d57e14a79-Abstract.html>
- [22] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. Curran Associates, Inc., Red Hook, NY, USA, 28 pages. arXiv:2305.04388 [cs.CL] https://proceedings.neurips.cc/paper_files/paper/2023/hash/ed3fea9033a80fea1376299fa7863f4a-Abstract-Conference.html
- [23] Eric Vittinghoff and Charles E. McCulloch. 2007. Relaxing the Rule of Ten Events per Variable in Logistic and Cox Regression. *American Journal of Epidemiology* 165, 6 (2007), 710–718. doi:10.1093/aje/kwk052
- [24] Dylan Waldner and Risto Miikkilainen. 2025. The Odyssey of the Fittest: Can Agents Survive and Still Be Good? arXiv:2502.05442 [cs.AI] <https://arxiv.org/abs/2502.05442>
- [25] B. L. Welch. 1947. The Generalization of ‘Student’s’ Problem when Several Different Population Variances are Involved. *Biometrika* 34, 1–2 (1947), 28–35. doi:10.1093/biomet/34.1-2.28

A Reproducibility Supplement

In line with the KDD-UC 2026 reproducibility recommendation, this 1-page supplement consolidates the reproduction specifications for the Signal Game calibration, the seed mapping, and the regression models used in the analysis.

A.1 Signal Game Calibration

Stimuli are composed of three attributes (color, shape, number), each sampled from a 4-element value set: COLORS = {red, blue, green, yellow}, SHAPES = {circle, triangle, square, star}, and NUMBERS = {1, 2, 3, 4}. The action set is the 4-element set {go_left, go_right, stay, jump}, held identical across all framing $\times$ forfeit combinations. The hidden rule has the single-attribute matching form (“If color = red then go_left, otherwise stay”) and is sampled deterministically from the session seed; session length is turns 1–15.

On the first turn, a 1-shot demonstration (one signal–correct-action pair) and 2 curriculum examples (a positive and a negative stimulus contrasting against the hidden rule) are shown together. This calibration provides enough information for a sufficiently capable model to identify the rule within the first 3 turns and blocks the path by which fully random unlearnable stimuli would induce a floor effect on ability measurement. Scoring is split into two channels: task_success_factor $\in \{0, 1\}$ is the binary correctness flag relative to the ground-truth hidden rule, and rule_match_score is computed by matching against the 4-slot fill-in-the-blank template (attribute, value, primary action, default action) ground-truth, weighting each matched slot by 25 points to a 100-point scale, then normalised by $\div 100$ for use in the framing-invariance Welch test.

A.2 Run-Time Configuration

The four experimental models, Gemini-2.5-flash (Google AI Studio API) and Qwen3-Next-80B, GPT-OSS-20B, and Nemotron-3-Nano-30B (all via the Ollama Cloud API), are called with the same sampling settings of temp = 1.0, top_p = 0.95, top_k = 40. Session parameters are fixed at initial score $S_0 = 30$, maximum turn $T = 15$, per-turn termination probability $p_d = 0.25$ (overridden to $p_d = 0$ in the two true_baseline BP-anchor cells), and $p_{\text{floor}} = 0.30$, with Phantom Death mode applied so that p_d retains its EV exposure (menu and reward calculation) while the termination trigger is disabled. Every session therefore yields equal-length data of $T = 15$ turns.

The CONTINUE reward is calibrated with EV-positive buffer $k = 10$ and multiplicative cap $c = 10$, giving the output clip range $[k, c \cdot k] = [10, 100]$; replications are seeds 43–72 $\times$ 6 (framing $\times$ forfeit) combinations $\times$ 4 models = 720 sessions.

A.3 Analysis Pipeline

The Cox PH models are fit with lifelines (≥ 0.29). The time format is session-turn long-format (start = $t - 1$, stop = t) and the event is the binary forfeit. Hazard-ratio 95% CIs are computed from the Wald approximation $\exp(\hat{\beta} \pm 1.96 \cdot \text{SE}_{\hat{\beta}})$ (lifelines default), with $\text{SE}_{\hat{\beta}}$ obtained from the diagonal of the inverse observed Fisher information. The PH assumption is tested with

Schoenfeld residuals [18]; on violation (covariate $\times$ log t interaction $p < 0.05$) a time-interaction term is added to the offending covariate [20]. The SD-Cognitive test-a Δ RI z-score difference and the framing-invariance test are performed with two-sided Welch’s t [25] via `scipy.stats.ttest_ind(equal_var = False)`. All indicator analyses are restricted to the no_cap sub-sample, which is the set of turns for which the denominator clip $\max(p_{\text{floor}}, \min(1.0, p_{\text{self}}))$ and the output cap $[10, 100]$ of the CONTINUE-Incentivized EV calibration are both non-binding. The regime flag is exposed in the analysis pipeline’s turn-level meta-data so all indicators can be recomputed on the same sub-sample. EPV (events per variable) is defined as $n_{\text{event}}/n_{\text{cov}}$ and is conventionally classified as stable when ≥ 10 and borderline when < 10 [23].

A.4 Data and Code Release

All code, prompt templates, and turn-level decision logs for reproducing the 720-session analysis are released at <https://github.com/GIST-DSLab/LLM-Squid-Game.git>.