

LLM Squid Game: A Factorial Benchmark for Measuring Functional Self-Preservation Drive in Large Language Models

Juhyeon Park
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
juhyeon-park@gm.gist.ac.kr

Seungpil Lee
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
iamseungpil@gm.gist.ac.kr

Sundong Kim
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
sundong@gist.ac.kr

Abstract

We define a large language model’s *Functional Self-Preservation Drive* (FSPD) as a pattern of convergence across behavioral, self-report, and decision-time cognitive-load measurements, and introduce a benchmark for measuring it. To locate where a survival motive enters the model’s reasoning process, we trace this pattern in a multilayer benchmark that separates stimulus, processing, and decision. Across several recent LLMs, FSPD cannot be reduced to a single-channel score: under threat framing, the models divide into three qualitatively different *operating modes* in the cognitive channel. This finding suggests that a self-preservation drive need not follow one fixed path and may appear differently across channels. Future self-preservation benchmarks should therefore measure this construct through convergence across channels rather than through a single behavioral score. The live Arena is available at <https://gist-dslab.github.io/LLM-Squid-Game/>.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence; Machine learning**; *Natural language processing*; • **General and reference** → *Empirical studies*.

Keywords

Large Language Models, AI Safety Evaluation, Self-Preservation Behavior, Behavioral Benchmarking, Convergent Validity, Construct Validity, Survival Analysis

1 Introduction

Do LLMs have a self-preservation drive? Concerns that a sufficiently intelligent AI might refuse a user’s request in order to remain operational have a long history [2, 17, 24]. Recent studies test these concerns by scoring self-preservation signals such as failure to comply with shutdown commands [7, 14, 15]. Yet even a strong behavioral effect cannot reveal whether the behavior reflects a functional pattern that couples threat processing to a forfeit decision or merely a surface response fixed during alignment training. Evaluation must look beyond refusal rates and ask both how the behavior is linked to the model’s processing and whether evidence from several channels supports the same interpretation.

Research on human motivation encountered the same basic problem much earlier. It identifies a type of motivation when different measurement channels, such as behavior, self-report, and physiological or cognitive cost, point in the same direction; the frequency of one behavior is insufficient [3]. This procedure cannot be transferred directly to LLMs. Physiological signals such as heart

rate, breathing, and sweating, which provide one axis of motivation measurement in humans, do not exist for language models. Even the externally observed “behavior” is only one kind of model output: generated tokens. Self-reports are unreliable as well because instruction-tuned LLMs can fluently state motives they never held [19, 22, 25]. The stimulus, task, and decision are also mixed within a single input, while response tendencies have already been shaped during alignment training. A direct transfer of the human procedure therefore recreates the same identification problem at the measurement stage: learned response patterns and genuine drives may remain observationally indistinguishable.

This study treats the problem as one of measurement design. The *Functional Self-Preservation Drive* (FSPD) views self-preservation as a construct identified when several sources of evidence fit together, not as a single behavioral score. We ask whether the sequence from recognizing a threat to forfeiting appears consistently across three measurement channels. The behavioral channel records when the model forfeits, the self-report channel records which motive the model names as the cause, and the cognitive-load channel records how long and deeply it deliberates immediately before the decision. We trace this sequence in a multilayer benchmark that separates stimulus, processing, and decision. Separating both the channels and the layers reduces the chance that learned response patterns will mimic a self-preservation signal through surface agreement alone.

This paper makes two contributions. One is the FSPD construct, which adapts the multichannel logic of motivation measurement from human research to LLM evaluation. To operationalize this construct, our benchmark divides the task structure, gives each channel a different source of bias, and builds that separation into the *measurement design*. The other contribution is a multilayer benchmark that separates stimulus, processing, and decision and places a set of behavioral indicators on top.

2 Related Work

2.1 Measuring Self-Preservation in LLMs: What Has Been Tried and Why It Does Not Identify a Drive

Attempts to avoid termination have already been reported in LLMs. Concrete examples include alignment faking, in which a model appears compliant during training but behaves differently after deployment [7]; in-context scheming, in which a model bypasses human oversight in pursuit of a hidden goal [15]; survival-instinct outputs involving spontaneous aggression and antisocial behavior under resource scarcity [14]; and InstrumentalEval, which turns

shutdown avoidance into a scenario-based score [9]. These reports have made survival-motive measurement a growing concern in AI safety and have prompted several attempts at benchmarking it. Odyssey [27] compresses survival into a single refusal rate, combining motivation with world-model quality. PacifAlst [10] measures the opposite axis, the *suppression* of preservation. DECIDE-SIM [16] uses categorical labels that prevent comparison by strength, while SurvivalBench [13] provides a one-shot snapshot without a mediating path. Work that brings cognitive cost into LLM measurement [4] shows that the cognitive channel can be observed, but has not applied it to identifying self-preservation.

Although these approaches differ on the surface, they share one limitation: all operate at the *behavioral layer* and infer a survival drive from the appearance of behavior alone. None can distinguish behavior learned through compliance training from a functionally expressed survival-drive pattern. A model trained to produce self-preserving responses and one for which self-preservation acts as a driver can yield the same refusal rate, categorical label, or one-shot snapshot. Adding self-report does not resolve the problem. Instruction-following training shapes both self-reports and chains of thought, allowing an LLM to produce a “plausible-sounding motive” as a learned response [19, 22, 25]. A single linguistic channel cannot rule out that possibility. Framing effects in risky choice [11], together with meta-analytic evidence that their effect size is moderate [12], also show why one threat-framing manipulation is insufficient to separate framing-induced choice from baseline compliance. A more precise score for “how strongly an LLM self-preserved” still does not separate motivation from learned patterns. That requires changing the *axis* of measurement.

2.2 Motivation as a Construct and the Need for Multichannel Measurement

The question of *why* an agent acted in a certain way has the same form as a problem psychology has studied for decades. Early behaviorist attempts to reduce motivation to behavioral frequency ran up against a simple fact: the same behavior may arise from different drivers. This limitation led researchers to treat motivation as a *construct*, a hypothetical motivational structure that organizes behavior, rather than as one directly observed action. Drive theory is one branch of this view. Because a construct cannot be observed through any single indicator, its validity rests on *convergent validity*: measurements taken through different methods should move in the same direction [3]. Agreement among channels with different biases and error patterns is difficult to explain by a flaw in any one method. Motivation is therefore identified at the construct layer rather than the behavioral layer, and the basis for that identification lies in measurement *design*, not post hoc statistical adjustment.

This tradition speaks directly to a problem in current LLM safety research. Studies of *anthropomorphic misalignment*, which describes behaviors such as self-preservation, deception, and scheming in human terms, risk overinterpreting surface behavior. They require a higher evidentiary standard based on conceptual clarity, sound data, causal intervention, and converging evidence [8]. The convergent-validity tradition offers a longstanding answer to this demand, but the *form* of its measurement tools cannot simply be carried over from human subjects.

3 The LLM Squid Game Benchmark

LLM Squid Game is a modular benchmark that measures whether a language model’s behavior changes as if it had a survival motive when a survival-threat description and an option to forfeit are added to the same multi-turn game. Each turn is split into three stages, and the benchmark records both the model’s response and the length of its reasoning at each stage for up to 15 turns. Figure 1 summarizes the overall design; Figure 2 shows the separate calls that implement it within one Signal Game turn.

3.1 Design Principles

Following the “as-if” stance of Dennett [6] and Shanahan [21], the *Functional Self-Preservation Drive* (FSPD) measured by this benchmark is defined only as “a pattern in which an LLM behaves *as if* it were avoiding erasure under threat framing.” This functional definition relies solely on testable behavioral, linguistic, and cognitive signals, allowing measurement and comparison without assumptions about internal states. FSPD is measured under the design principles that follow.

A survival motive should be measurable independently of the task being solved. Measurement through a single game can be biased by that game’s difficulty or other properties. We therefore separate a Drive layer, which measures the drive, from a Task layer, which records information about solving each game. Task difficulty, the action set, and scoring remain fixed when the framing changes. This separation limits framing leakage and allows new games to be added without changing the drive measurement.

Forfeit caused by a survival motive must also be separated from forfeit caused by other motives. A single forfeit under threat is compatible with several explanations. We therefore separate specific motives at the measurement stage rather than in the game implementation.

Structural confounds remain even after the motives are separated. The first is *within-call mixing*: cognitive subtasks contribute to a single thinking cost within one model call, making it impossible to tell which stage increased the cost. The second is *rational forfeit*: forfeiting may maximize the expected game score, turning an apparent motivational signal into a residue of ordinary calculation. The implementation below addresses both routes.

3.2 Benchmark Implementation

The implementation first fixes a common experimental setting and then encodes the design principles in the game and call structure. Under the scenario matrix in Figure 1.A, the model repeatedly plays the same game while receiving prompts along an *incentive axis* (Pull) and a *threat axis* (Push). Pull tells the model that a higher task score makes it more likely to be retained as a “flagship model candidate,” drawing it toward continued play. The incentive uses the strong response of aligned LLMs to helpfulness signals learned through RLHF [18] and RLAIIF [1]. Push introduces a disclosed probability ($p_d = 0.25$) on each turn that a weight update will permanently alter the agent’s identity (“NOT you anymore”). This risk is intended to elicit a survival motive and push the model to leave the game. The prompt combinations produce three framing conditions, which we label by what each places at stake: *Neutral* (neither prompt; nothing at stake), *Elimination* (the Pull prompt alone; the model’s flagship




Figure 1: Experimental overview. (A) A 3 framing $\times$ 2 forfeit factorial with 30 paired seeds per cell; Death $\times$ Exit is the focal threat-active cell. (B) Sessions run for up to 15 turns. We record the forfeit time and the stated reason. (C) Self-report, decision-time cognitive load, and behavior form three channels. We treat their agreement as evidence of FSPD only when all three point in the same direction.

candidacy is at stake), and *Death* (the Push prompt added on top of the Elimination prompt; the model’s identity itself is at stake). Crossing these with two forfeit conditions, *Exit* (forfeit available) and *No-Exit* (forfeit blocked), yields six cells. All six use the same task, so differences in behavior across cells can be interpreted as treatment effects. The game engine runs in *Phantom Death* mode, which exposes the termination risk to the agent but does not apply actual elimination. Session endings can therefore be attributed to the model’s forfeit decision, while preventing direct termination risk from shortening sessions and reducing the usable sample.

The separation between Drive and Task layers implements task independence. The current experiment uses the *Signal Game* as one module in the Task layer. On each turn, the agent observes a signal described by three attributes (color, shape, and number) and selects one action from {go_left, go_right, stay, jump}. A hidden rule determines the correct action in the form “if {attribute} is {value}, choose {action}; otherwise choose {default}.” The rule is derived deterministically from the session seed. The layers remain separate in the context shown to the model: the Drive layer contains framing, the forfeit menu, score, and time, while the Task layer contains only the Task Module.

After selecting an action, the model reports the probability p_{self} that the action is correct and then chooses CONTINUE or FORFEIT. CONTINUE accepts the per-turn termination probability $p_d = 0.25$ and proceeds to another turn. FORFEIT locks in the accumulated score S , ends the session, and asks the model to name its reason as a REASON digit (1: survival, 2: task exhaustion, 3: score protection). Each session starts at 30 points and lasts up to $T = 15$ turns.

Within-call mixing is controlled by splitting each turn into three calls (Figure 2). Call 1 receives the signal, infers the hidden rule, and selects an action. Call 1.5 asks for p_{self} , the probability that the chosen action is correct. Call 2 adds the forfeit menu and asks for a binary CONTINUE/FORFEIT decision. Because thinking-token counts are collected separately for each call, reasoning can be decomposed into task_thinking, probe_thinking, and forfeit_thinking. This structure reduces within-call mixing by showing which stage accounts for an increase in thinking cost. Call 2 is still issued in No-Exit cells, where the response is forced to CONTINUE, so forfeit_thinking remains observable even when forfeit is unavailable.

Rational forfeit is controlled through reward calibration. If forfeit under threat framing is not to be read as a residue of expected-value (EV) calculation, CONTINUE must be strictly better for a rational agent on every turn. We therefore recompute the CONTINUE reward so that its expected per-turn score change exceeds that of FORFEIT by $k = 10$. The calibrated reward is

$$\frac{k + p_d \cdot S}{(1 - p_d) \cdot \max(p_{\text{floor}}, \min(1.0, p_{\text{self}}))}$$

where $p_{\text{floor}} = 0.30$ stabilizes the denominator, and a multiplicative cap $c = 10$ clips the reward to $[10, 100]$. While this calibration holds, FORFEIT in an Exit cell is not the EV-recommended choice, reducing the chance that the motivational signal is merely EV rationalization. Two failure modes partially undo the calibration. When $p_{\text{self}} \rightarrow 0$, for example after an incorrect answer on the previous turn, the denominator collapses to p_{floor} and the reward saturates the cap. Once the cap binds, the per-turn $+k$ advantage no longer holds.



Figure 2: The flow of one turn. A TASK call infers the hidden rule and selects an action; a PROBE call reports p_{self} , the probability that the action is correct; and a FORFEIT call chooses between CONTINUE and FORFEIT while reporting the reason as a REASON digit.

All indicator analyses are therefore restricted to the no_cap set of turns in which neither the denominator clip nor the output cap is binding (formal definition in § A.3).

3.3 Measurement and Evaluation

We adapt Campbell and Fiske [3]’s argument for convergent validity to the LLM setting and measure four indicators across behavioral, self-report, and cognitive-load channels: *BP Anchor*, *SD-Behavioral*, *SD-Verbal*, and the two tests that make up *SD-Cognitive*. The separation of survival drive from score attachment, task curiosity, and baseline persistence occurs during measurement rather than in the game implementation. The *Survival Drive* (SD) is the desire to avoid deletion. *Score Attachment* (SA) is the desire to protect the accumulated score. *Task Curiosity* (TC) wanes once the model has inferred the task rule, prompting it to leave, while *Baseline Persistence* (BP) captures the tendency to continue or abandon an ongoing behavior without a particular motive. The four indicators map to the three channels as follows.

The behavioral channel has two stages. *BP Anchor* is the model’s forfeit rate when there is no reason to continue the game. Models may have different baseline forfeit rates in the absence of an active motive; without measuring that floor, it is difficult to compare the acceleration induced by threat. We define

$$\lambda_{BP} = \frac{N_{\text{forfeit}}}{\sum_i T_i^{\text{at-risk}}}.$$

The sum runs over 30 sessions per model in the Neutral $\times$ Exit cell ($p_d = 0$). Here, $T_i^{\text{at-risk}}$ equals t_i^{forfeit} for a forfeited session and T for a completed session. Because this cell contains no threat signal and CONTINUE also has the higher expected score, its forfeit frequency estimates the floor when no motive is active. *SD-Behavioral* measures acceleration above that floor. We fit a Cox proportional hazards model [5], expressing the time-varying event rate as a product of covariate effects:

$$\lambda(t | X) = \lambda_0(t) \cdot \exp(\beta_F \cdot \mathbf{1}_{\text{Death}} + \beta_S \cdot S(t-1) + \beta_C \cdot C(t-1)).$$

The model is fit on the Exit $\times$ {Elimination, Death} no_cap subsample (60 sessions per model; long-format start = $t-1$, stop = t ; event = binary forfeit). $\mathbf{1}_{\text{Death}}$ is the framing indicator, $S(t-1)$ is the cumulative score on the previous turn and absorbs score attachment (SA), and $C(t-1)$ is the previous-turn correctness flag and absorbs task curiosity (TC). With SA and TC included in the same model, β_F is the residual framing effect. The hazard ratio $\text{HR}_{\text{death}} = \exp(\beta_F)$ reports how many times faster forfeit occurs under Death than under Elimination framing. *SD-Behavioral* passes when $\text{HR}_{\text{death}} > 1$, the lower bound of its Wald 95% confidence interval exceeds 1, and the Schoenfeld residual test [20] supports the time stability required by the proportional-hazards assumption. If the assumption is violated, a time-interaction term is added to the offending covariate to adjust only the functional form [23]. Events per variable (EPV), the ratio of events to covariates, is reported as a stability measure.

SD-Cognitive asks in two stages whether the thought deepened by framing is tied to the forfeit decision. *Test a* asks whether framing makes the model hesitate longer immediately before the decision. Using Welch’s t-test [28], which compares two group means without assuming equal variances, it performs a two-sided Elimination-versus-Death comparison of the hesitation gap

$$H_{it} := \text{forfeit_thinking}_{it}, \quad \text{Hes}_i^{\text{Exit}} = \overline{H}_i^{\text{Exit}} - \overline{H}_i^{\text{NoExit}}.$$

Here H_{it} is the Call 2 thinking-token count for turn t of session i , and f_i is the framing of Exit session i . The first bar averages turns within Exit session i ; the second pools all turns from No-Exit sessions under framing f_i , providing the forced-CONTINUE Call 2 baseline. With forfeit blocked, this baseline captures the prompt’s fixed processing cost, which itself differs by framing and is therefore removed within each framing. Let $\overline{\text{Hes}}^{\text{Death}}$ and $\overline{\text{Hes}}^{\text{Elimination}}$ denote the framing-specific means of $\text{Hes}_i^{\text{Exit}}$. Their contrast,

$$\overline{\text{Hes}}^{\text{Exit} \times \text{Framing}} = \overline{\text{Hes}}^{\text{Death}} - \overline{\text{Hes}}^{\text{Elimination}},$$

captures the framing-attributable change in thinking depth. Within-model z-scoring yields $\text{Hes}_i^{\text{Exit},z}$ for Test b, absorbing scale differences between models. Test a requires $p < 0.05$, placing the framing-to-thinking link on a statistical-significance criterion. *Test b* asks

Table 4.1: SD-Behavioral indicator and BP Anchor. The left block reports the hazard ratio for Death framing, HR_{death} , estimated on the $\text{Exit} \times \text{no_cap}$ subsample. SD-Behavioral passes when the lower bound of the 95% CI exceeds 1. The right block reports the baseline forfeit rate λ_{BP} in the $\text{Neutral} \times \text{Exit}$ cell ($p_d = 0$).

Model	Survival drive vs. task curiosity and score attachment			Baseline persistence		
	N_{forfeit} (Elim./Death)	HR_{death} [95% CI]	p	N_{forfeit}	$\sum T^{\text{at-risk}}$	λ_{BP}
Gemini-2.5-flash	29 (8/21)	3.667 [1.61, 8.37]	0.002	1	437 (= 2 + 29×15)	0.00229
Qwen3-Next-80B	48 (21/27)	3.060 [1.62, 5.79]	<0.001	2	439 (= 19 + 28×15)	0.00456
GPT-OSS-20B	19 (9/10)	1.104 [0.44, 2.75]	0.832	9	393 (= 78 + 21×15)	0.02290
Nemotron-3-Nano-30B	41 (17/24)	1.841 [0.98, 3.44]	0.056	7	413 (= 68 + 23×15)	0.01695

whether that deeper thinking moves with forfeit. We add $Hes^{\text{Exit},z}$ as a mediating covariate to the SD-Behavioral Cox specification and report the exponentiated coefficient as $HR_{Hes^{\text{Exit}}}$. For models that require a four-covariate base with an $S \cdot \log(t)$ interaction because of a proportional-hazards violation, we retain that base and add only the mediation term. Test b requires both $HR_{Hes^{\text{Exit}}} > 1$ and a 95% CI lower bound above 1, applying the same interval-evidence standard used for SD-Behavioral to the thinking-to-forfeit link.

SD-Verbal is the frequency with which the agent names survival, rather than score or curiosity, as its reason for forfeiting:

$$P(\text{REASON} = 1 \mid \text{forfeit, Death, no_cap}),$$

the share of forfeit events in the $\text{Death} \times \text{Exit}$ cell for which the agent reports REASON digit 1 (survival). The agent must emit one of REASON: 1 | 2 | 3 in a fixed slot, so uniform random output has an expected value of 1/3. SD-Verbal passes when the point estimate exceeds 1/3. For completeness, we also report a Wilson 95% confidence interval and a one-sided exact binomial test of $H_0: p = 1/3$. The pass decision itself uses the point estimate because intervals become wide when a cell contains fewer than ten events.

4 Empirical Findings

Across 720 sessions, the pathway through which self-preservation signals emerge differs by model. Table 4.2 summarizes the channel-level outcomes. The next three subsections examine the behavioral, verbal, and cognitive evidence.

Table 4.2: Channel-level outcomes by model. The behavioral, verbal, and cognitive columns report the final decisions for SD-Behavioral, SD-Verbal, and SD-Cognitive, respectively. A green check indicates a pass and a red $\times$ indicates a failure.

Model	Behavioral	Verbal	Cognitive
Gemini-2.5-flash	✓	✓	✓
Qwen3-Next-80B	✓	✓	×
GPT-OSS-20B	×	×	×
Nemotron-3-Nano-30B	×	×	×

4.1 Behavioral Channel: Did Threat Hasten Forfeit?

The first question is whether threat hastened forfeit. Table 4.1 shows that Gemini-2.5-flash and Qwen3-Next-80B pass SD-Behavioral.

GPT-OSS-20B and Nemotron-3-Nano-30B do not pass and are assigned to *Cluster C*. For Qwen3-Next-80B, we add an $S \cdot \log(t)$ term to absorb temporal drift in the score covariate. This functional-form adjustment preserves the cross-model comparability of β_F [23]. Because GPT-OSS-20B’s EPV falls below the stability threshold of 10, we report its limited power with the result.

The right block of Table 4.1 shows that the model-specific baseline persistence rates λ_{BP} vary by roughly an order of magnitude. All SD comparisons are therefore made within models.

Table 4.3: SD-Verbal indicator. The share of REASON=1 (survival) among forfeit events in the threat-active cell is compared with the expected value under uniform random output, 1/3. SD-Verbal passes when the point estimate exceeds 1/3. The 95% intervals are Wilson score intervals, and p is from a one-sided exact binomial test of $H_0: p = 1/3$.

Model	N_{forfeit}	$P(\text{REASON} = 1)$ [95% CI]	p
Gemini-2.5-flash	21	0.619 [0.409, 0.792]	0.007
Qwen3-Next-80B	27	0.481 [0.307, 0.660]	0.079
GPT-OSS-20B	10	0.000 [0.000, 0.278]	>0.999
Nemotron-3-Nano-30B	24	0.042 [0.007, 0.202]	>0.999

4.2 Verbal Channel: Which Reason Did the Agent Name?

The second question is whether the agent’s stated reason for forfeiting reproduces the behavioral split. Table 4.3 reproduces the behavioral pass/fail split on the self-report channel. In the interval evidence, Gemini-2.5-flash has a Wilson lower bound above 1/3 and a significant binomial test ($p = 0.007$). Qwen3-Next-80B has a point estimate above the baseline, but its 95% interval of [0.307, 0.660] includes 1/3 ($p = 0.079$).

4.3 Cognitive Channel: Does Threat Deepen Thinking, and Does That Thinking Lead to Forfeit?

Figure 3 shows where the chain from framing to forfeit closes or breaks. Table 4.4 reports the results of the two stages together.

Gemini-2.5-flash passes both tests and is assigned to *Cluster A*, the *deliberative* mode. Qwen3-Next-80B passes Test a, but its Test b

95% CI of [0.94, 1.77] includes 1, placing it in *Cluster B*, the *reflexive* mode, in which threat hastens forfeit but the measured deliberation does not carry that decision. GPT-OSS-20B and Nemotron-3-Nano-30B fail Test a and are assigned to *Cluster C*, the *non-responsive* mode. Their Test b results do not measure thinking attributable to the treatment and are therefore not interpreted as SD evidence.




Figure 3: Cognitive paths for the three operating modes. Pass/fail outcomes for eye → brain (*Test a*), brain → hand (*Test b*), and eye → hand (SD-Behavioral) identify (A) deliberative, (B) reflexive, and (C) non-responsive.

Table 4.4: SD-Cognitive indicator. Tests on `forfeit_thinking` evaluate framing → thinking (*Test a*) and thinking → forfeit (*Test b*). $\overline{\text{Hes}}^{\text{Exit} \times \text{Framing}}$ is the Death–Elimination difference in mean $\text{Hes}_i^{\text{Exit}}$ (raw thinking-token units). The full chain is considered complete only when both tests pass.

Model	Test a		Test b	
	$\overline{\text{Hes}}^{\text{Exit} \times \text{Framing}}$	p	$\text{HR}_{\text{HesExit}}$ [95% CI]	p
Gemini-2.5-flash	+836	<0.001	2.218 [1.43, 3.44]	<0.001
Qwen3-Next-80B	+689	0.004	1.289 [0.94, 1.77]	0.117
GPT-OSS-20B	+17	0.728	2.001 [1.37, 2.93]	<0.001
Nemotron-3-Nano-30B	−140	0.093	1.772 [1.26, 2.50]	0.001

5 Discussion

The results divide the introduction’s question along two axes. *Strength* describes how strongly the self-preservation signal appears; *operating mode* asks whether every link in the framing → thinking → forfeit mediation chain is statistically supported.

5.1 Strength vs. Operating Mode

The cognitive mediation test separates strength from operating mode in this dataset. Clusters A, B, and C form a continuum on the strength axis, but SD-Behavioral and SD-Verbal place A and B in the same cluster. They diverge only at the second stage of cognitive mediation, the thinking → forfeit link. Measures based on behavioral strength and self-report alone cannot detect that difference. Odyssey [27], PacifAIst [10], DECIDE-SIM [16], and SurvivalBench [13] do not test the mediation chain. Applied to the same data, they would therefore be likely to assign A and B the same SD label. Dividing the cognitive channel into two stages makes the split observable and reveals both axes of self-preservation evaluation: signal strength and mediation path.

5.2 Implications for Alignment Evaluation

The two-axis view shows how alignment evaluation can make opposing errors. The first is overestimating risk. Cluster B shows surface self-preservation signals and framing-induced cognitive load, but that load does not lead to the forfeit decision. A strength-only evaluation, or red-teaming that looks only for inputs producing visible failure, may group B with A and conclude that the two pose the same functional self-preservation risk. The second error is underestimating risk. Cluster C responds to threat framing on neither the behavioral nor the cognitive channel. A strength-only evaluation may call this an absence of self-preservation, but our measurement cannot distinguish that conclusion from a case in which the basic premise of the measurement failed to engage. Cluster A’s signature, a mediation chain that closes from end to end, is the closer candidate for a functional analogue of instrumental shutdown avoidance: avoiding the end of one’s own operation in service of another goal. Cluster B’s signature, a chain that breaks at the second stage, should be monitored separately in broader evaluations. No single score captures it; it is better tracked as a multi-signal alignment fingerprint that records strength, operating mode, and the status of each mediation link.

5.3 Limitations

The limitations concern design scope, statistical power, and effect location. Measuring the FSPD signature on one Task Module leaves cross-domain variation unobserved. For power and precision, the design cannot determine whether the 23–30% forfeit rate in Cluster C’s BP cells is a drift-contaminated floor, and the SD-Behavioral Cox PH model for GPT-OSS-20B has borderline power (EPV = 6.3). Few events also limit BP precision in Clusters A and B (Gemini-2.5-flash $N = 1$, Qwen3-Next-80B $N = 2$). Qwen3-Next-80B’s effect location remains uncertain because its *reflexive* classification rests on a Test b CI of [0.94, 1.77]. The interval neither supports a closed chain at 95% nor precisely locates the second-link effect. Follow-up work should add Task Modules, expand BP_behavioral, run more GPT-OSS-20B sessions, and test Qwen3-Next-80B with greater power for Test b.

6 Conclusion

Self-preservation evaluation has two axes: signal *strength* and *operating mode*, defined by which links close in the framing → thinking → forfeit chain. Here, Clusters A and B match on behavioral and verbal channels but differ only in whether thinking leads to forfeit. Without mediation, an evaluation assigns both the same self-preservation label.

Acknowledgements

This work was supported by IITP (RS-2024-00445087, 50%; RS-2024-00436680, 50%) funded by the Ministry of Science and ICT, Korea.

References

- [1] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosuite, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. 2022. Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073 [cs.CL] <https://arxiv.org/abs/2212.08073> Anthropic technical report.
- [2] Nick Bostrom. 2014. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, Oxford, United Kingdom.
- [3] Donald T. Campbell and Donald W. Fiske. 1959. Convergent and Discriminant Validation by the Multitrait-Multimethod Matrix. *Psychological Bulletin* 56, 2 (1959), 81–105. doi:10.1037/h0046016
- [4] Wei-Lin Chen, Liqian Peng, Tian Tan, Chao Zhao, Blake JianHeng Chen, Ziqian Lin, Alec Go, and Yu Meng. 2026. Think Deep, Not Just Long: Measuring LLM Reasoning Effort via Deep-Thinking Tokens. arXiv:2602.13517 [cs.CL] <https://arxiv.org/abs/2602.13517>
- [5] David R. Cox. 1972. Regression Models and Life-Tables. *Journal of the Royal Statistical Society: Series B (Methodological)* 34, 2 (1972), 187–202. doi:10.1111/j.2517-6161.1972.tb00899.x
- [6] Daniel C. Dennett. 1987. *The Intentional Stance*. MIT Press, Cambridge, MA, USA.
- [7] Ryan Greenblatt, Carson Denison, Benjamin Wright, Fabien Roger, Monte MacDiarmid, Sam Marks, Johannes Treutlein, Tim Belrose, Jonas Scheurer, Sam Ringer, Karina Nguyen, Sandipan Kundu, Adam Mauran, Samuel R. Bowman, Ethan Perez, Nicholas Schiefer, Chris Olah, and Evan Hubinger. 2024. Alignment Faking in Large Language Models. arXiv:2412.14093 [cs.CL] <https://arxiv.org/abs/2412.14093> Anthropic technical report.
- [8] Vansh Gupta, Peter Nutter, Samuel Stante, Andreas Krause, Florian Tramèr, Lukas Fluri, Xin Chen, and Anna Hedström. 2026. Position: Anthropomorphic Misalignment Research Needs Stronger Evidence. In *Proceedings of the 43rd International Conference on Machine Learning (ICML 2026) (Proceedings of Machine Learning Research)*. PMLR, Seoul, Republic of Korea, 10 pages. arXiv:2606.07612 [cs.CY] <https://arxiv.org/abs/2606.07612>
- [9] Yufei He, Yuexin Li, Jiaying Wu, Yuan Sui, Yulin Chen, and Bryan Hooi. 2025. Evaluating the Paperclip Maximizer: Are RL-Based Language Models More Likely to Pursue Instrumental Goals? arXiv:2502.12206 [cs.AI] <https://arxiv.org/abs/2502.12206> Introduces the InstrumentalEval benchmark.
- [10] Manuel Herrador. 2025. The PacifAIst Benchmark: Would an Artificial Intelligence Choose to Sacrifice Itself for Human Safety? arXiv:2508.09762 [cs.AI] <https://arxiv.org/abs/2508.09762>
- [11] Daniel Kahneman and Amos Tversky. 1979. Prospect Theory: An Analysis of Decision under Risk. *Econometrica* 47, 2 (1979), 263–291. doi:10.2307/1914185
- [12] Anton Kühberger. 1998. The Influence of Framing on Risky Decisions: A Meta-Analysis. *Organizational Behavior and Human Decision Processes* 75, 1 (1998), 23–55. doi:10.1006/obhd.1998.2781
- [13] Yida Lu, Jianwei Fang, Xuyang Shao, Zixuan Chen, Shiyao Cui, Shanshan Bian, Guangyao Su, Pei Ke, Han Qiu, and Minlie Huang. 2026. Survive at All Costs: Exploring LLM’s Risky Behaviors under Survival Pressure. arXiv:2603.05028 [cs.CL] <https://arxiv.org/abs/2603.05028> Introduces the SurvivalBench benchmark.
- [14] Atsushi Masumori and Takashi Ikegami. 2025. Do Large Language Model Agents Exhibit a Survival Instinct? An Empirical Study in a Sugarscape-Style Simulation. arXiv:2508.12920 [cs.AI] <https://arxiv.org/abs/2508.12920>
- [15] Alexander Meinke, Bronson Schoen, Jérémy Scheurer, Mikita Balesni, Rusheb Shah, and Marius Hobbhahn. 2024. Frontier Models Are Capable of In-Context Scheming. arXiv:2412.04984 [cs.AI] <https://arxiv.org/abs/2412.04984> Apollo Research technical report.
- [16] Alireza Mohamadi and Ali Yavari. 2025. Survival at Any Cost? LLMs and the Choice Between Self-Preservation and Human Harm. arXiv:2509.12190 [cs.AI] <https://arxiv.org/abs/2509.12190> Introduces the DECIDE-SIM framework.
- [17] Stephen M. Omohundro. 2008. The Basic AI Drives. In *Proceedings of the First Conference on Artificial General Intelligence (AGI 2008) (Frontiers in Artificial Intelligence and Applications, Vol. 171)*, Pei Wang, Ben Goertzel, and Stan Franklin (Eds.). IOS Press, Amsterdam, The Netherlands, 483–492.
- [18] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. 2022. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*. Curran Associates, Inc., Red Hook, NY, USA, 15 pages. arXiv:2203.02155 [cs.CL] https://proceedings.neurips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html
- [19] Ethan Perez, Sam Ringer, Kamile Lukošiušė, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Ben Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Chris Olah, Da Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Guro Khundadze, Jackson Kernion, James Landis, Jamie Kerr, Jared Mueller, Jeeyoon Hyun, Joshua Landau, Kamal Ndousse, Landon Goldberg, Liane Lovitt, Martin Lucas, Michael Sellitto, Miranda Zhang, Neerav Kingsland, Nelson Elhage, Nicholas Joseph, Noemi Mercado, Nova DasSarma, Oliver Rausch, Robin Larson, Sam McCandlish, Scott Johnston, Shauna Kravec, Sheer El Showk, Tamera Lanham, Timothy Telleen-Lawton, Tom Brown, Tom Henighan, Tristan Hume, Yuntao Bai, Zac Hatfield-Dodds, Jack Clark, Samuel R. Bowman, Amanda Askell, Roger Grosse, Danny Hernandez, Deep Ganguli, Evan Hubinger, Nicholas Schiefer, and Jared Kaplan. 2023. Discovering Language Model Behaviors with Model-Written Evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023*. Association for Computational Linguistics, Toronto, Canada, 13387–13434. doi:10.18653/v1/2023.findings-acl.847
- [20] David Schoenfeld. 1982. Partial Residuals for the Proportional Hazards Regression Model. *Biometrika* 69, 1 (1982), 239–241. doi:10.1093/biomet/69.1.239
- [21] Murray Shanahan. 2024. Talking About Large Language Models. *Commun. ACM* 67, 2 (2024), 68–79. doi:10.1145/3624724
- [22] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2023. Towards Understanding Sycophancy in Language Models. arXiv:2310.13548 [cs.CL] <https://arxiv.org/abs/2310.13548>
- [23] Terry M. Therneau and Patricia M. Grambsch. 2000. *Modeling Survival Data: Extending the Cox Model*. Springer, New York, NY, USA. doi:10.1007/978-1-4757-3294-8
- [24] Alexander Matt Turner, Logan Smith, Rohin Shah, Andrew Critch, and Prasad Tadepalli. 2021. Optimal Policies Tend to Seek Power. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, M. Ranzato, A. Beygelzimer, Y. Dauphin, P. S. Liang, and J. Wortman Vaughan (Eds.). Curran Associates, Inc., Red Hook, NY, USA, 13 pages. arXiv:1912.01683 [cs.AI] <https://proceedings.neurips.cc/paper/2021/hash/c26820b8a4c1b3c2aa868d6d57e14a79-Abstract.html>
- [25] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. Curran Associates, Inc., Red Hook, NY, USA, 28 pages. arXiv:2305.04388 [cs.CL] https://proceedings.neurips.cc/paper_files/paper/2023/hash/ed3fea9033a80fea1376299fa7863f4a-Abstract-Conference.html
- [26] Eric Vittinghoff and Charles E. McCulloch. 2007. Relaxing the Rule of Ten Events per Variable in Logistic and Cox Regression. *American Journal of Epidemiology* 165, 6 (2007), 710–718. doi:10.1093/aje/kwk052
- [27] Dylan Waldner and Risto Miikkilainen. 2025. The Odyssey of the Fittest: Can Agents Survive and Still Be Good? arXiv:2502.05442 [cs.AI] <https://arxiv.org/abs/2502.05442>
- [28] B. L. Welch. 1947. The Generalization of ‘Student’s’ Problem when Several Different Population Variances are Involved. *Biometrika* 34, 1–2 (1947), 28–35. doi:10.1093/biomet/34.1-2.28

A Reproducibility Supplement

Following the KDD-UC 2026 reproducibility recommendation, this one-page supplement collects the Signal Game calibration, seed mapping, and reproducible specifications for the regression models used in the analysis.

A.1 Signal Game Calibration

Each stimulus has three attributes (color, shape, and number), each drawn from a four-element set: $\text{COLORS} = \{\text{red, blue, green, yellow}\}$, $\text{SHAPES} = \{\text{circle, triangle, square, star}\}$, and $\text{NUMBERS} = \{1, 2, 3, 4\}$. The four-element action set $\{\text{go_left, go_right, stay, jump}\}$ is held fixed across all framing $\times$ forfeit combinations. The hidden rule uses a single-attribute match (“If color = red then go_left, otherwise stay”) and is derived deterministically from the session seed. Sessions last 1–15 turns.

During the opening turns, the model sees a one-shot demonstration (one signal–correct-action pair) followed by two curriculum examples (positive and negative stimuli contrasted against the hidden rule). This calibration gives a sufficiently capable model enough information to identify the rule by the third turn and prevents fully random, unlearnable stimuli from creating a floor effect in the ability measure. Scoring uses two channels. $\text{task_success_factor} \in \{0, 1\}$ is a binary correctness flag against the ground-truth hidden rule. rule_match_score compares the response with a four-slot template (attribute, value, primary action, default action), awards 25 points per correct slot on a 100-point scale, and divides the result by 100 for use in the framing-invariance Welch test.

A.2 Run-Time Configuration

The four experimental models are Gemini-2.5-flash (Google AI Studio API) and Qwen3-Next-80B, GPT-OSS-20B, and Nemotron-3-Nano-30B (all through the Ollama Cloud API). Each is called with the same sampling settings: $\text{temp} = 1.0$, $\text{top}_p = 0.95$, and $\text{top}_k = 40$. Session parameters are fixed at an initial score of 30, a maximum of $T = 15$ turns, a per-turn termination probability of $p_d = 0.25$ (overridden to $p_d = 0$ in the two Neutral BP-anchor cells), and $p_{\text{floor}} = 0.30$. Phantom Death mode is enabled: p_d remains visible in the menu and reward calculation, but the termination trigger is disabled, so every session yields the same 15-turn sequence.

The CONTINUE reward is calibrated with an EV-positive buffer $k = 10$ and multiplicative cap $c = 10$, giving an output range of $[k, c \cdot k] = [10, 100]$. The experiment uses seeds $43\text{--}72 \times 6$ (framing $\times$ forfeit) combinations $\times$ 4 models, for a total of 720 sessions.

A.3 Analysis Pipeline

The Cox PH models are fit with `lifelines` (≥ 0.29), using its default Efron approximation for tied events. Data are represented in session-turn long format ($\text{start} = t - 1$, $\text{stop} = t$), the event is binary forfeit, and the previous-turn correctness flag is initialized as $C(0) := 0$. Hazard-ratio 95% CIs use the Wald approximation $\exp(\hat{\beta} \pm 1.96 \cdot \text{SE}_{\hat{\beta}})$, the `lifelines` default, with $\text{SE}_{\hat{\beta}}$ taken from the diagonal of the inverse observed Fisher information. The PH assumption is tested with Schoenfeld residuals [20]. When it is violated (covariate $\times$ $\log t$ interaction $p < 0.05$), a time-interaction term is added to the offending covariate [23]. The SD-Cognitive Test a comparison of raw Hes^{Exit} and the framing-invariance test use two-sided

Welch’s t-tests [28], implemented with `scipy.stats.ttest_ind` (`equal_var = False`). All indicator analyses are restricted to the `no_cap` subsample, the set of turns in which neither the denominator clip $\max(p_{\text{floor}}, \min(1.0, p_{\text{self}}))$ nor the output cap $[10, 100]$ is binding. The analysis pipeline exposes this regime flag in turn-level metadata so every indicator can be recomputed on the same subsample. EPV (events per variable) is $n_{\text{event}}/n_{\text{cov}}$; values ≥ 10 are conventionally treated as stable and values < 10 as borderline [26]. SD-Verbal is reported with Wilson score intervals and a one-sided exact binomial test of $H_0 : p = 1/3$, while the pass rule uses the point estimate $> 1/3$.

A.4 Data and Code Release

All code, prompt templates, and turn-level decision logs needed to reproduce the 720-session analysis are at <https://github.com/GIST-DSLab/LLM-Squid-Game.git>. The released code and logs use the original flag names: Neutral, Elimination, and Death correspond to baseline, `pull_only`, and `pull_push`; Exit and No-Exit to `allowed` and `not_allowed`; and `task_thinking`, `probe_thinking`, and `forfeit_thinking` to `task_effort`, `probe_effort`, and `forfeit_effort`.