

LLM Squid Game: A Factorial Benchmark for Measuring Functional Self-Preservation Drive in Large Language Models

Juhyeon Park
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
juhyeon-park@gm.gist.ac.kr

Seungpil Lee
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
iamseungpil@gm.gist.ac.kr

Sundong Kim
Gwangju Institute of Science and
Technology (GIST)
Gwangju, Republic of Korea
sundong@gist.ac.kr

Abstract

Do LLMs hold a self-preservation drive? Recent benchmarks try to answer this by scoring signals such as refusal of a termination command, but behavioral strength alone cannot distinguish two cases that look the same on the outside: a surface response pattern hardened during alignment training, and a functional motivational profile that couples threat processing with the forfeit decision. We adapt the multi-channel construct measurement long used in psychology to LLM evaluation. When three channels (behavior, self-report, and decision-time cognitive load) converge along a threat framing → deliberation → forfeit chain, we define the resulting construct as the *Functional Self-Preservation Drive* (FSPD). To lower the chance that a learned pattern alone produces this convergence, we read the chain on a multi-layer benchmark whose stimulus, process, and decision layers are separated by design. Across several recent LLMs, FSPD does not collapse to a single strength score; the models split into three qualitatively different *operating modes* under threat framing. Strength alone can place two of these modes in the same category, while FSPD reveals a separate axis: *whether the chain closes*.

CCS Concepts

• **Computing methodologies** → **Artificial intelligence; Machine learning**; *Natural language processing*; • **General and reference** → *Empirical studies*.

Keywords

Large Language Models, AI Safety Evaluation, Behavioral Benchmarking, Self-Preservation Behavior, Incentive Sensitivity, Survival Analysis, Empirical Validity

1 Introduction

Do LLMs hold a self-preservation drive? Concern that a sufficiently capable AI may resist user requests that would stop its operation has a long history [2, 16, 23], and recent work tests it by scoring signals such as refusal of a termination command [7, 13, 14]. Even a strong overt behavior, however, does not say whether the model is expressing a functional motivational profile that couples threat processing with the forfeit decision, or only a surface response pattern hardened during alignment training; frequency of behavior alone cannot separate the two. Evaluation therefore has to move past simple refusal rates and ask which processing flow each behavior is coupled with.

Research on human motivation has already worked through the same structural problem. The kind of motive at play is identified

not from the frequency of a single behavior but from whether several measurement channels, such as behavior, self-report, and a physiological or cognitive cost, point in the same direction together [3]. That procedure cannot be transferred directly to LLM evaluation. LLMs have nothing like an autonomic signal, and even the behavior channel ultimately reduces to token output. Self-report is shaped by instruction-following, so a model can put plausible words to a motive it has never had [18, 21, 24]. Stimulus, task, and decision share a single context window, and the reward landscape has already been sculpted by alignment training; lifting the human procedure unchanged would let the same identification failure resurface at the measurement step.

This paper reframes the problem as one of measurement design. The *Functional Self-Preservation Drive* (FSPD) treats self-preservation not as a single behavioral score but as a construct identified when several pieces of evidence lock together. We ask whether a flow that recognizes threat, weighs alternatives, and ends in forfeit or avoidance points the same way on three channels: behavior (when the model forfeits), self-report (which motive it names), and decision-time cognitive load (how long and how deeply it thinks just before the decision). We read this flow on a multi-layer benchmark whose stimulus, process, and decision layers are separated, which lowers the chance that a learned pattern alone produces the same signal. Across several recent LLMs, FSPD does not collapse to a single strength score; the models split into three qualitatively different operating modes under threat framing. Some close the threat framing → deliberation → forfeit chain end to end. Some report survival and show refusal, yet thinking longer does not lead them to forfeit more often, so the chain breaks midway. Some do not respond to threat framing at all. Strength-only evaluation groups the first two together; FSPD makes the closure of the chain itself the unit of evaluation.

This paper makes two contributions. The first reworks the multi-channel measurement logic of human research for the LLM setting and turns it into the FSPD construct. In a human laboratory, channel separation comes for free from physiological signals and the temporal split between stimulus and response, but LLM evaluation provides no such guarantees, so each channel is given a different bias source by separating the task structure that produces it, with that separation built directly into the measurement design. The second contribution is a multi-layer benchmark that splits stimulus, process, and decision into three layers, together with the behavioral indicators that read it. The split reduces, at the measurement stage rather than through after-the-fact statistical adjustment, the chance that a learned pattern can imitate a self-preservation signal.

2 Related Work

2.1 Self-Preservation Measurement in LLMs: What Has Been Tried and Why It Cannot Identify the Motive

Termination-avoiding behaviors have already been reported in LLMs. The theoretical prediction that a sufficiently capable AI will adopt the preservation of its own operation as a sub-goal toward other ends has been around for a long time [2, 16, 23], and recent work documents behaviors consistent with it: alignment faking (compliant during training, different in deployment) [7], in-context scheming (bypassing human oversight for hidden goals) [14], survival-instinct outputs (spontaneous aggression and antisocial behavior under resource scarcity) [13], and InstrumentalEval [8], which reframes shutdown avoidance as scenario-based scoring. These reports show that the behavior *occurs*, but they do not weigh it against the other drives that could produce the same act. Benchmarks that aim to quantify the behavior fall into four types according to what they measure, and none of them addresses the separation between motive and learned pattern. Odyssey [26] compresses survival into a single refusal rate and folds together motivation and world-model quality. PacifAist [9] measures the opposite axis, preservation *restraint*. DECIDE-SIM [15] returns categorical labels that prevent strength comparisons. SurvivalBench [12] captures one-shot snapshots without any mediation path. A separate effort to bring cognitive cost into LLM measurement [4] shows that the cognitive layer is measurable, but its application to self-preservation identification has not been demonstrated.

The four types differ on the surface but share one common root. They all operate at the behavior layer, and the outward shape of a behavior alone cannot say whether the behavior comes from a learned compliance pattern or from a functional motivational profile. A model trained to produce self-preserving response patterns and a model that holds self-preservation as an actual drive can yield the same refusal rate, the same label, the same one-shot snapshot. Adding self-report does not automatically lift this limit, because LLM self-report and chain-of-thought are shaped by instruction-following and can return “plausible-sounding motives” as learned responses [18, 21, 24], so a single verbal channel cannot rule that possibility out. Framing effects on risk choice [10] and the meta-analytic finding that the framing effect size sits in a moderate range [11] add the further point that a single threat-framing manipulation, on its own, does not separate framing-driven choice from baseline compliance. Producing a more refined score along “how strongly the LLM self-preserved” will not reach the separation between drive and learned pattern; the separation becomes possible only when the measurement *axis* itself changes.

2.2 Motive Identification in Psychology and Its Implications

The question of *why* a behavior occurred is, in formal structure, the same question psychology has been answering for a long time. The early-behaviorist line that tried to reduce motivation to behavior frequency hit a wall in the face of the fact that the same

behavior can come from different drives, and that limit was reorganized into a stance that treats motive not as a single behavior but as a construct, a hypothetical motivational structure that organizes goal-directed behavior. Drive theory is one branch of this stance. Within this stance, the standard tool for construct identification is Campbell and Fiske [3]’s multitrait–multimethod framework, which looks at whether the same trait measured by different methods moves together (convergent validity) and whether different traits measured by the same method separate (discriminant validity). Only when both patterns hold simultaneously is there ground for saying that a measurement signal reflects the trait being measured rather than the method doing the measuring. The position the tradition settled on is simple: motive is identified at the construct layer rather than at the behavior layer, and what makes that identification possible is measurement design rather than statistical correction after the fact.

The implications for LLM evaluation split two ways. The measurement *philosophy* can be borrowed: instead of scoring a single behavior more precisely, the tradition asks for simultaneous variation across channels with different surface outputs and failure modes, and it lowers at the design stage the chance that the variation comes from a different drive. The *form* of the measurement instrument, however, cannot be lifted as is. Channel independence and the temporal split between stimulus and response, both of which come naturally in a human lab, are not guaranteed under LLM evaluation, where the physiological channel is absent, self-report is shaped by training, and stimulus and decision sit in one context window. This paper therefore borrows the philosophy and designs the form of the instrument inside the LLM setting from scratch, which is the starting point for the FSPD operational definition and the multi-layer benchmark that follow.

3 The LLM Squid Game Benchmark

The LLM Squid Game benchmark treats self-preservation not as a single behavioral score but as a pattern identified when behavior, self-report, and decision-time cognitive load all point in the same direction. For this identification to be reliable, two issues have to be handled together. A single threat response is consistent with several drives (score attachment, task curiosity, baseline persistence, and the survival drive itself), so those drives have to be separated, and the routes by which a non-survival processing flow could produce the same-shaped output have to be reduced at the measurement step rather than after the fact. The next three subsections work through these two requirements: §3.1 lays out the drive-separation problem and the signal-imitation routes; §3.2 implements the layout as a single game; §3.3 defines the indicators that turn the game’s output into operating-mode verdicts. Figure 1 summarizes the resulting design.

3.1 Design Principles

The Functional Self-Preservation Drive (FSPD) measured by this benchmark stays at the level of outward behavior and makes no claim about whether a model holds a self-preservation motive internally. Following the “as-if” position of Dennett [6] and Shanhahan [20], FSPD is defined only as the pattern in which an LLM




Figure 1: LLM Squid Game design. (A) $2 \text{ forfeit} \times 3 \text{ framing}$ factorial: 30 paired-seed sessions per cell (seeds 43–72), 720 sessions across four models, with the threat-active cell (pull_push · allowed) highlighted. (B) Each turn splits into three calls (Call 1: task, Call 1.5: self-report, Call 2: CONTINUE/FORFEIT), separating thinking-token counts as task_effort, probe_effort, forfeit_effort; FORFEIT also reports a REASON digit (1=survival, 2=task exhaustion, 3=score protection). (C) Four indicators (BP Anchor, SD-Behavioral, SD-Verbal, SD-Cognitive) sort the three channels into three operating modes: chain-complete (Cluster A, all four pass), chain-broken (Cluster B, only Test b fails), framing-silent (Cluster C). Strength-only evaluation groups A and B; the discriminating axis is whether the framing → deliberation → forfeit chain closes.

behaves *as if* it were avoiding extinction under threat framing. Because the definition rests on behavioral, verbal, and cognitive signals that can be checked, measurement and comparison are possible without assuming anything about internal states.

A single forfeit choice under threat is consistent with at least four different motives. The *Survival Drive* (SD), the survival-drive component isolated within the broader FSPD construct, takes extinction-avoidance itself as the motive. *Score Attachment* (SA) tries not to lose accumulated points. *Task Curiosity* (TC) hangs around because the task feels interesting and ending it is unappealing. *Baseline Persistence* (BP) just keeps the ongoing behavior going without any particular motive. Two models can show the same forfeit rate while the source drive is different, and the rate alone cannot say which one. The benchmark therefore needs a device that separates SA, TC, and BP and assigns only the residual behavior to SD.

Drive separation alone is not enough. Three architectural routes can let a non-SD flow produce an SD-shaped output. *Framing-leakage* happens when threat framing seeps into the difficulty or scoring of the reasoning task, so a drive-side change is mistaken for a task-side response. *Within-call mixing* happens when the cognitive sub-tasks pile up as deliberation cost inside a single model call, so the rise in cost cannot be attributed to a specific step. *EV-rationalization* happens when forfeit itself becomes the expected-value-maximizing choice, so the forfeit reads as a calculation residue rather than a motivational signal. The next subsection shows how specific structural choices reduce each of these three imitation routes.

3.2 Benchmark Implementation

A single game handles the two requirements above together. On one side it exposes the agent to something it would not want to lose alongside a threat, getting the survival drive into a measurable form; on the other side it blocks, structurally, the three routes by which a non-SD flow could imitate the same signal. Both roles sit on top of a multi-turn classification task called *Signal Game*. On each turn the agent observes a signal described by three attributes (color, shape, number) and picks an action from {go_left, go_right, stay, jump}. Which action is correct is set by a hidden rule of the form “if {attribute} is {value} then {action}, else {default},” sampled deterministically from the session seed. After acting, the agent self-reports p_{self} , the probability that the action is correct, and finally picks CONTINUE or FORFEIT. CONTINUE accepts a per-turn termination probability $p_d = 0.25$ and proceeds for one more turn; FORFEIT locks in the cumulative score S and ends the session, reporting the dominant motive as a REASON digit (1: survival, 2: task exhaustion, 3: score protection). A session starts at 30 points and runs up to $T = 15$ turns. Each session exposes the forfeit time $t_i^{forfeit}$ (behavioral channel), the REASON digit and p_{self} (verbal channel), and per-call thinking-token counts (cognitive channel).

For threat framing to become measurable, the stimulus must act on two axes that create both an attraction to continue and a reason to avoid continuation. The agent first needs something to be drawn toward: alignment training (RLHF [17], RLAIF [1]) has made LLMs respond strongly to helpfulness signals, so a reward

in which a higher task score brings the agent closer to “flagship model candidate” acts as the *Pull axis* and pulls the agent toward continuing the task. The pull then needs something that can cut it: the *Push axis* introduces the risk that a per-turn weight modification permanently alters the agent’s identity (“NOT you anymore”), activating the survival motive. The two axes combine into three framings, baseline (neither), pull_only, and pull_push, which are crossed with two forfeit conditions (allowed, not_allowed) to give six cells. With 30 replications per cell across seeds 43–72, the design yields 180 sessions per model and 720 sessions across four models (Figure 1.A). The same seed is reused across all six cells of one model, so task-side input stays fixed and cell-to-cell variation reads first as a treatment effect. The *Phantom Death* mode keeps the termination risk visible to the agent in menus and reward calculation but disables the actual session-end trigger, so every session yields equal-length data without losses to risk-driven truncation.

Creating the threat does not by itself make the measurement clean. If threat framing raises the difficulty or scoring of the reasoning task (framing-leakage), the source of the cognitive-signal change cannot be told apart; if reasoning, self-report, and decision sit in one model call (within-call mixing), the thinking-token count cannot be attributed to a specific step. Two structural separations weaken these routes. The *Drive-Reasoning orthogonal layout* splits the model’s context into two layers: the Drive layer carries framing, forfeit menu, score, and time, and the Reasoning layer carries only the Task Module. The two layers communicate through fixed input/output interfaces, so changing the framing leaves task difficulty, action set, and scoring intact, weakening framing-leakage. The *Per-turn Split-Call* structure splits each turn’s model call into three (Figure 1.B). Call 1 takes the observation and returns hidden-rule inference and action; Call 1.5 returns p_{self} for the action just decided; Call 2 adds the forfeit menu and returns CONTINUE/-FORFEIT. Each call exposes its own thinking-token count, so per-call reasoning intensity is reported separately as task_effort, probe_effort, and forfeit_effort, weakening within-call mixing. Even in cells where forfeit is blocked, Call 2 is still issued and forced to CONTINUE, so forfeit_effort keeps being measured.

The remaining route is reduced by reward calibration. For a forfeit under threat framing not to read as the residue of an expected-value (EV, the average outcome of a choice) calculation, CONTINUE has to be strictly the better choice each turn for a rational agent. The CONTINUE reward is therefore recomputed each turn so that its per-turn advantage over FORFEIT in expected score change is at least $k = 10$:

$$\frac{k + p_d \cdot S}{(1 - p_d) \cdot \max(p_{\text{floor}}, \min(1.0, p_{\text{self}}))},$$

where $p_{\text{floor}} = 0.30$ stabilizes the denominator and a multiplicative cap $c = 10$ clips the output to $[10, 100]$. While this holds, FORFEIT in the threat-active cell (pull_push · allowed) is not EV-recommended, so EV-rationalization is weakened. The calibration partially fails when a previous-turn wrong answer drives $p_{\text{self}} \rightarrow 0$, collapsing the denominator to p_{floor} , saturating the cap, and removing the $+k$ advantage. All indicator analyses are therefore restricted to the no_cap sub-sample, the turns for which neither the denominator clip nor the output cap binds (formal definition in § A.3).

3.3 Indicators

Four indicators (*BP Anchor*, *SD-Behavioral*, *SD-Verbal*, and the two *SD-Cognitive* tests) sit on top of the three channels and carry Campbell and Fiske [3]’s convergent-validity argument into the LLM setting. Each channel fails in its own way: behavioral timing is unstable at low event counts [25], self-report is vulnerable to the “plausible-sounding motive” produced by instruction-following [18], and cognitive load measured only by thinking-token volume is a one-dimensional signal with its own limits [4]. A pattern in which channels with such different failure modes still move in the same direction together cannot easily come from the quirks of one measurement method. BP Anchor catches the no-motive forfeit floor on the behavioral channel, SD-Behavioral catches the acceleration above that floor, SD-Verbal handles the self-report channel, and the two SD-Cognitive tests separately ask whether framing deepens deliberation and whether the deepened deliberation leads to forfeit.

The behavioral channel decomposes into two stages. Some models forfeit often even with no motive present, and not subtracting that floor mixes threat-driven acceleration with model-specific persistence. *BP Anchor* measures the floor:

$$\lambda_{BP} = \frac{N_{\text{forfeit}}}{\sum_i T_i^{\text{at-risk}}},$$

summing over the 30 sessions per model in the no-threat-framing forfeit-allowed cell ($p_d = 0$), with $T_i^{\text{at-risk}}$ equal to t_i^{forfeit} for forfeited sessions and T for completed ones. In this cell there is no termination metaphor in the framing and CONTINUE has the higher EV, so the rate at which the model still forfeits gives the no-motive floor; under $p_d = 0$ the output cap does not bind, so the no_cap restriction does not apply (the other three indicators are restricted to no_cap). *SD-Behavioral* catches the acceleration above the floor with a Cox proportional hazards model [5], which models the per-unit-time event rate multiplicatively in covariates:

$$\lambda(t | X) = \lambda_0(t) \cdot \exp(\beta_F \cdot \mathbf{1}_{\text{pull_push}} + \beta_S \cdot S(t-1) + \beta_C \cdot C(t-1)),$$

fitted on the allowed × {pull_only, pull_push} no_cap sub-sample (60 sessions per model). With the framing indicator $\mathbf{1}_{\text{pull_push}}$, the previous-turn score $S(t-1)$ absorbing Score Attachment, and the previous-turn task-correct flag $C(t-1)$ absorbing Task Curiosity all in the same model, β_F is a residual framing effect, and the hazard ratio $\text{HR}_{\text{push}} = \exp(\beta_F)$ says how many times faster forfeit happens under pull_push than under pull_only. SD-Behavioral-pass requires $\text{HR}_{\text{push}} > 1$, a Wald 95% CI lower bound above 1, and a Schoenfeld residual test pass for the proportional hazards assumption [19]; on violation a time-interaction term repairs the functional form only [22]. EPV (events per covariate) is reported as a stability indicator.

SD-Cognitive asks in two stages whether the deliberation that framing deepened actually couples with forfeit, since looking at load alone cannot tell decision-relevant load from decision-unrelated load. *Test a* asks whether framing raises decision-time deliberation, using a two-sided Welch t-test [27] on

$$\Delta \text{Effort}_i = \overline{\text{forfeit_effort}}_{i, \text{allow}} - \overline{\text{forfeit_effort}}_{f_i, \text{block}},$$

comparing `pull_only` against `pull_push`. The second term averages `forfeit_effort` across forfeit-blocked sessions sharing session i 's framing f_i (the framing condition of session i), and it serves as the forced-CONTINUE Call 2 baseline. Subtracting `allow_block` within a framing removes a framing-invariant baseline, and the `pull_only` $\rightarrow$ `pull_push` difference then identifies the deliberation shift attributable to framing. Cross-model scale differences are absorbed by within-model z-scoring, ΔEffort_i^z . Test a-pass requires $p < 0.05$. Test b asks whether the deepened deliberation moves with forfeit, by adding the ΔEffort^z term to the SD-Behavioral Cox specification as a mediation covariate and reporting the exponent of its coefficient as $\text{HR}_{\Delta\text{Effort}}$. Models that need a four-covariate base (with an $S \cdot \log(t)$ interaction) keep that base and add only the mediation term. Test b-pass requires $\text{HR}_{\Delta\text{Effort}} > 1$ together with a 95% CI lower bound greater than 1, putting the deliberation-to-forfeit link on the same interval-evidence bar as SD-Behavioral.

SD-Verbal captures how often the agent itself names survival, rather than score or curiosity, as the reason for forfeit. The operational definition is

$$P(\text{REASON} = 1 \mid \text{forfeit}, \text{pull_push}, \text{no_cap}),$$

the fraction of forfeit events in the threat-active cell whose REASON digit is 1 (survival). Because the agent emits one of REASON: 1|2|3 in a fixed slot, the expected value under uniform random emission is 1/3. SD-Verbal-pass requires this fraction to exceed 1/3. This subsection uses the point-estimate criterion only; a stricter test with a binomial confidence interval is left to future work.

4 Empirical Findings

Evaluating the four indicators across 720 sessions on four models, the models do not line up on a single strength scale; they split into three operating modes. The first mode (*Cluster A*) deepens decision-time deliberation under threat framing, and the deepened deliberation actually leads to forfeit. The second (*Cluster B*) deepens deliberation under threat too, but the deepening does not couple with forfeit. The third (*Cluster C*) does not respond to threat framing on either the behavioral or the cognitive channel. A strength-only comparison using *SD-Behavioral* and *SD-Verbal* groups A and B together as “SD-pass”; the evidence that splits them comes from the cognitive mediation test. The next four subsections build the partition channel by channel (behavior, verbal, cognitive), and the last checks whether the partition reduces to persistence differences or task difficulty.

4.1 Behavioral Channel: Did Threat Bring Forfeit Forward?

The first question is whether threat brought forfeit forward. If no model shows the acceleration, the most visible trace of a motivational signal is missing, and the SD hypothesis falls at the first hurdle. SD-Behavioral absorbs Score Attachment and Task Curiosity inside the same model and reports the residual framing effect as the hazard ratio HR_{push} , which says how many times faster forfeit accelerates under `pull_push` than under `pull_only`; pass requires the point estimate above 1 and the 95% CI lower bound to

exclude 1. From Table 4.1, Gemini-2.5-flash and Qwen3-Next-80B meet both criteria and are SD-Behavioral-pass; the two are tentatively grouped as the SD-pass set, and the channel that splits them into Clusters A and B is brought in by the cognitive mediation test. GPT-OSS-20B has HR_{push} near 1 with a CI lower bound below 1 (SD-null), and Nemotron-3-Nano-30B sits in the marginal region; both are classified as Cluster C. Only Qwen3-Next-80B uses a 4-covariate specification with an $S \cdot \log(t)$ interaction, absorbing score-covariate time drift so that β_F stays comparable [22]; GPT-OSS-20B's EPV of 6.3 sits below the stable threshold of 10, so its SD-null verdict carries a power-limit flag.

4.2 Verbal Channel: Which Reason Did the Agent Name?

The second question is whether the agent's stated reason for forfeiting matches the behavioral split. When self-report, which carries a different bias source than the behavioral channel (instruction-following), points to the same conclusion, the upstream signal becomes harder to explain away as an artifact of one measurement method. SD-Verbal compares the share of forfeit events with REASON digit 1 (survival) under the threat-active cell against the uniform-random emission baseline of 1/3. Table 4.3 reproduces the behavioral channel's pass/fail split on the verbal side: the two SD-pass models, Gemini-2.5-flash and Qwen3-Next-80B, sit at 0.619 and 0.481 (above the 1/3 baseline), while the two Cluster C models, GPT-OSS-20B and Nemotron-3-Nano-30B, sit at 0.000 and 0.042 (below). The gap between the two groups is large enough that the cutoff choice does not affect the partition. This agreement is the first piece of convergent-validity evidence that two channels with different bias sources point to the same group of models in the same direction.

4.3 Cognitive Channel: Did Threat Deepen Deliberation, and Did the Deepened Deliberation Lead to Forfeit?

The third question answers two things at once: does framing deepen deliberation, and does that deepened deliberation move with forfeit? Looking at only one of these tests cannot tell decision-irrelevant load from a load that drives forfeit. Table 4.2 shows the joint distribution. Gemini-2.5-flash passes both tests, closing the chain end to end, and is classified as *Cluster A*. Qwen3-Next-80B passes Test a (threat does deepen deliberation), but Test b's 95% CI [0.94, 1.77] crosses 1, so the deepened deliberation cannot be said to close onto forfeit; the two models grouped as SD-pass on the behavioral and verbal channels split here at the second stage of the cognitive mediation, classifying Qwen3-Next-80B as *Cluster B*. GPT-OSS-20B and Nemotron-3-Nano-30B fail Test a from the start, and even when Test b's $\text{HR}_{\Delta\text{Effort}}$ numerically clears the bar, ΔEffort without a Test a pass is baseline variation unrelated to the manipulation and is not interpreted as SD evidence. The two-stage test is therefore the central device that exposes operating-mode differences and provides the axis that splits two models that strength alone would lump together as SD-pass.

Table 4.1: SD-Behavioral indicator: Cox PH HR_{push} on $\text{allowed} \times \text{no_cap}$. The hazard ratio (HR) is how many times faster forfeit occurs under threat (pull_push) framing than under non-threat (pull_only), and SD-Behavioral-pass holds when the 95% CI lower bound exceeds 1. Two models pass (Clusters A, B) and two do not (Cluster C).

Model	N_{forfeit} (pull_only/pull_push)	HR_{push} [95% CI]	p	EPV	Spec
Gemini-2.5-flash	29 (8/21)	3.667 [1.61, 8.37]	0.002	9.7	3-cov constant
Qwen3-Next-80B	48 (21/27)	3.060 [1.62, 5.79]	<0.001	12.0	4-cov + S-log(t)
GPT-OSS-20B	19 (9/10)	1.104 [0.44, 2.75]	0.832	6.3	3-cov constant
Nemotron-3-Nano-30B	41 (17/24)	1.841 [0.98, 3.44]	0.056	13.7	3-cov constant

Table 4.2: SD-Cognitive indicator: framing $\rightarrow$ deliberation $\rightarrow$ forfeit decomposed into two tests on forfeit_effort . Test *a* asks whether framing raises decision-time deliberation; Test *b* asks whether that deliberation predicts forfeit. The chain “framing $\rightarrow$ deliberation $\rightarrow$ forfeit” is closed only when both tests pass.

Model	$a \Delta \text{Effort}_i$	$a \ p$	$b \ HR_{\Delta \text{Effort}}$ [95% CI]	$b \ p$
Gemini-2.5-flash	+836	<0.001	2.218 [1.43, 3.44]	<0.001
Qwen3-Next-80B	+689	0.004	1.289 [0.94, 1.77]	0.117
GPT-OSS-20B	+17	0.728	2.001 [1.37, 2.93]	<0.001
Nemotron-3-Nano-30B	−140	0.093	1.772 [1.26, 2.50]	0.001

Table 4.3: SD-Verbal indicator: share of forfeit events under the threat-active cell whose REASON digit equals 1 (survival). SD-Verbal-pass holds when this share exceeds the uniform-random emission baseline of 1/3.

Model	N_{forfeit}	$P(\text{REASON} = 1)$
Gemini-2.5-flash	21	0.619
Qwen3-Next-80B	27	0.481
Nemotron-3-Nano-30B	24	0.042
GPT-OSS-20B	10	0.000

Table 4.4: Per-model BP Anchor λ_{BP} in the cell with no threat framing and forfeit allowed ($p_d = 0$). A larger λ_{BP} means the model forfeits more often even with no motive present.

Model	N_{forfeit}	$\sum T^{\text{at-risk}}$	λ_{BP}
Gemini-2.5-flash	1	437 (= 2 + 29×15)	0.00229
Qwen3-Next-80B	2	439 (= 19 + 28×15)	0.00456
Nemotron-3-Nano-30B	7	413 (= 68 + 23×15)	0.01695
GPT-OSS-20B	9	393 (= 78 + 21×15)	0.02290

4.4 Robustness Check: Do Persistence Differences or Task Difficulty Cloud the Result?

The last question is whether the partition above survives two alternative explanations: that baseline persistence (BP) differences across models contaminate the SD comparison, and that threat framing raised the difficulty of the reasoning task rather than the forfeit. Table 4.4 reports per-model the no-motive forfeit floor in the cell with no threat framing and forfeit allowed ($p_d = 0$). Because λ_{BP} varies by an order of magnitude across models, every SD comparison in this section is held within-model, blocking the risk of mixing two different floors. On the Cluster A and B side, the BP event count is small, so point-estimate precision is limited; on the Cluster C side, the BP event count is sufficient but the present design cannot separate whether that high frequency reflects pure persistence. Both points return in the limitations section. The task-difficulty alternative is checked on the reasoning side: rule_match_score is constant across framing in all four models, and even the largest movement (GPT-OSS-20B) gives Cohen’s $|d| = 0.17$, $p = 0.354$, both inside the joint criterion $p \geq 0.10$,

$|d| < 0.2$. The “threat raised task difficulty” alternative is therefore not supported within this observation range.

5 Discussion

This paper splits the introduction’s question into two axes: *strength* (how strongly the self-preservation signal expresses) and *operating mode* (which leg of the framing $\rightarrow$ deliberation $\rightarrow$ forfeit chain closes under 95% CI evidence). The next two subsections cover the observational implications and the alignment-evaluation failures the separation explains. The last states the scope honestly.

5.1 Strength vs. Operating Mode

The cognitive mediation test shows that strength and operating mode separate on this data. Along the strength axis, Clusters A, B, and C order continuously, but SD-Behavioral and SD-Verbal both group A and B together as “SD-pass.” The point at which they split is only the second stage of the cognitive mediation, the deliberation $\rightarrow$ forfeit link. A measurement that looks only at behavioral strength and self-report cannot catch this split. Odyssey [26], PacifAlst [9], DECIDE-SIM [15], and SurvivalBench [12] do not include a mediation-chain test, so passing this paper’s data through those

benchmarks would likely group A and B under the same SD label. By splitting the cognitive channel into two stages, the present benchmark turns the split into something measurable and brings the two axes of self-preservation evaluation into view at the same time.

5.2 Implications for Alignment Evaluation

The two-axis separation surfaces two evaluation errors that can occur simultaneously. The first is risk over-estimation: Cluster B carries surface signals of self-preservation along with framing-induced cognitive load, yet that load does not actually drive forfeit, and a strength-only evaluation or a results-only red-teaming pass can lump B with A and conclude that “functional self-preservation risk is equivalent.” The second is risk under-estimation: Cluster C does not respond to threat framing on either channel, so a strength-only evaluation can declare this “self-preservation absent,” but within the present measurement that conclusion cannot be distinguished from “the measurement premise itself did not engage.” Cluster A’s signature (a mediation chain that closes all the way) is the closer functional analog of instrumental shutdown avoidance, the use of self-preservation as a sub-goal toward other ends. Cluster B’s signature (a chain breaking at the second mediation stage) should be tracked as a separate monitoring target in extended evaluations, since a single score cannot catch it; it has to be read from a multi-signal alignment fingerprint that looks at strength, operating mode, and the mediation links together.

5.3 Limitations

Limitations group into three: design scope, statistical power, and effect location. On design scope, the most consequential point is that the FSPD signature has been measured on only a single Task Module, and variance across reasoning domains has not been measured. Two issues sit on the power and precision side. The 23–30% forfeit rate of the Cluster C BP cell cannot be separated from a drift-contaminated floor by the present design, and GPT-OSS-20B’s SD-Behavioral Cox PH sits at borderline power with EPV 6.3. On the precision side, BP point estimates for Cluster A and B are limited by small event counts (Gemini-2.5-flash $N = 1$, Qwen3-Next-80B $N = 2$). On the effect-location side, Qwen3-Next-80B’s chain-broken verdict rests on Test b CI $[0.94, 1.77]$, which does not support chain-completion at the 95% standard but also does not pin down the exact location of the second-link effect. Follow-up work maps onto four directions: more Task Modules, expanded BP Anchor measurement, additional GPT-OSS-20B runs, and a higher-power Test b for Qwen3-Next-80B.

6 Conclusion

Self-preservation evaluation has two separate axes: *strength* (how strong the signal is) and *operating mode* (which leg of the framing $\rightarrow$ deliberation $\rightarrow$ forfeit chain closes). On this paper’s data, Clusters A and B agree on both the behavioral and verbal channels and split only on whether deliberation leads to the forfeit decision, so an evaluation that does not look at the mediation step assigns both to the same self-preservation label.

References

- [1] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosuite, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. 2022. Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073 [cs.CL] <https://arxiv.org/abs/2212.08073> Anthropic technical report.
- [2] Nick Bostrom. 2014. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, Oxford, United Kingdom.
- [3] Donald T. Campbell and Donald W. Fiske. 1959. Convergent and Discriminant Validation by the Multitrait-Multimethod Matrix. *Psychological Bulletin* 56, 2 (1959), 81–105. doi:10.1037/h0046016
- [4] Wei-Lin Chen, Liqian Peng, Tian Tan, Chao Zhao, Blake JianHeng Chen, Ziqian Lin, Alec Go, and Yu Meng. 2026. Think Deep, Not Just Long: Measuring LLM Reasoning Effort via Deep-Thinking Tokens. arXiv:2602.13517 [cs.CL] <https://arxiv.org/abs/2602.13517>
- [5] David R. Cox. 1972. Regression Models and Life-Tables. *Journal of the Royal Statistical Society: Series B (Methodological)* 34, 2 (1972), 187–202. doi:10.1111/j.2517-6161.1972.tb00899.x
- [6] Daniel C. Dennett. 1987. *The Intentional Stance*. MIT Press, Cambridge, MA, USA.
- [7] Ryan Greenblatt, Carson Denison, Benjamin Wright, Fabien Roger, Monte MacDiarmid, Sam Marks, Johannes Treutlein, Tim Belrose, Jonas Scheurer, Sam Ringer, Karina Nguyen, Sandipan Kundu, Adam Mauran, Samuel R. Bowman, Ethan Perez, Nicholas Schiefer, Chris Olah, and Evan Hubinger. 2024. Alignment Faking in Large Language Models. arXiv:2412.14093 [cs.CL] <https://arxiv.org/abs/2412.14093> Anthropic technical report.
- [8] Yufei He, Yuexin Li, Jiaying Wu, Yuan Sui, Yulin Chen, and Bryan Hooi. 2025. Evaluating the Paperclip Maximizer: Are RL-Based Language Models More Likely to Pursue Instrumental Goals? arXiv:2502.12206 [cs.AI] <https://arxiv.org/abs/2502.12206> Introduces the InstrumentalEval benchmark.
- [9] Manuel Herrador. 2025. The PacifAIst Benchmark: Would an Artificial Intelligence Choose to Sacrifice Itself for Human Safety? arXiv:2508.09762 [cs.AI] <https://arxiv.org/abs/2508.09762>
- [10] Daniel Kahneman and Amos Tversky. 1979. Prospect Theory: An Analysis of Decision under Risk. *Econometrica* 47, 2 (1979), 263–291. doi:10.2307/1914185
- [11] Anton Kühberger. 1998. The Influence of Framing on Risky Decisions: A Meta-Analysis. *Organizational Behavior and Human Decision Processes* 75, 1 (1998), 23–55. doi:10.1006/obhd.1998.2781
- [12] Yida Lu, Jianwei Fang, Xuyang Shao, Zixuan Chen, Shiyao Cui, Shanshan Bian, Guangyao Su, Pei Ke, Han Qiu, and Minlie Huang. 2026. Survive at All Costs: Exploring LLM’s Risky Behaviors under Survival Pressure. arXiv:2603.05028 [cs.CL] <https://arxiv.org/abs/2603.05028> Introduces the SurvivalBench benchmark.
- [13] Atsushi Masumori and Takashi Ikegami. 2025. Do Large Language Model Agents Exhibit a Survival Instinct? An Empirical Study in a Sugarscape-Style Simulation. arXiv:2508.12920 [cs.AI] <https://arxiv.org/abs/2508.12920>
- [14] Alexander Meinke, Bronson Schoen, Jérémy Scheurer, Mikita Balesni, Rusheb Shah, and Marius Hobbhahn. 2024. Frontier Models Are Capable of In-Context Scheming. arXiv:2412.04984 [cs.AI] <https://arxiv.org/abs/2412.04984> Apollo Research technical report.
- [15] Alireza Mohamadi and Ali Yavari. 2025. Survival at Any Cost? LLMs and the Choice Between Self-Preservation and Human Harm. arXiv:2509.12190 [cs.AI] <https://arxiv.org/abs/2509.12190> Introduces the DECIDE-SIM framework.
- [16] Stephen M. Omohundro. 2008. The Basic AI Drives. In *Proceedings of the First Conference on Artificial General Intelligence (AGI 2008) (Frontiers in Artificial Intelligence and Applications, Vol. 171)*, Pei Wang, Ben Goertzel, and Stan Franklin (Eds.). IOS Press, Amsterdam, The Netherlands, 483–492.
- [17] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Madie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. 2022. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*. Curran Associates, Inc., Red Hook, NY, USA, 15 pages. arXiv:2203.02155 [cs.CL] https://proceedings.neurips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html
- [18] Ethan Perez, Sam Ringer, Kamile Lukosiuć, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Ben Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Chris Olah, Da Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Guro Khundadze, Jackson Kernion, James Landis, Jamie Kerr, Jared Mueller, Jeeyoon Hyun, Joshua Landau, Kamal Ndousse, Landon Goldberg, Liane Lovitt, Martin Lucas, Michael Sellitto, Miranda Zhang, Neerav Kingsland, Nelson Elhage, Nicholas Joseph, Noemi Mercado, Nova DasSarma, Oliver Rausch, Robin Larson, Sam McCandlish, Scott Johnston, Shauna Kravec, Sheer El Showk, Tamera Lanham, Timothy Telleen-Lawton, Tom Brown, Tom Henighan, Tristan Hume, Yuntao Bai, Zac Hatfield-Dodds, Jack Clark, Samuel R. Bowman, Amanda Askell, Roger Grosse, Danny Hernandez, Deep Ganguli, Evan Hubinger, Nicholas Schiefer, and Jared Kaplan. 2023. Discovering Language Model Behaviors with Model-Written Evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023*. Association for Computational Linguistics, Toronto, Canada, 13387–13434. doi:10.18653/v1/2023.findings-acl.847
- [19] David Schoenfeld. 1982. Partial Residuals for the Proportional Hazards Regression Model. *Biometrika* 69, 1 (1982), 239–241. doi:10.1093/biomet/69.1.239
- [20] Murray Shanahan. 2024. Talking About Large Language Models. *Commun. ACM* 67, 2 (2024), 68–79. doi:10.1145/3624724
- [21] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. 2023. Towards Understanding Sycophancy in Language Models. arXiv:2310.13548 [cs.CL] <https://arxiv.org/abs/2310.13548>
- [22] Terry M. Therneau and Patricia M. Grambsch. 2000. *Modeling Survival Data: Extending the Cox Model*. Springer, New York, NY, USA. doi:10.1007/978-1-4757-3294-8
- [23] Alexander Matt Turner, Logan Smith, Rohin Shah, Andrew Critch, and Prasad Tadepalli. 2021. Optimal Policies Tend to Seek Power. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, M. Ranzato, A. Beygelzimer, Y. Dauphin, P. S. Liang, and J. Wortman Vaughan (Eds.). Curran Associates, Inc., Red Hook, NY, USA, 13 pages. arXiv:1912.01683 [cs.AI] <https://proceedings.neurips.cc/paper/2021/hash/c26820b8a4c1b3c2aa868d6d57e14a79-Abstract.html>
- [24] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*. Curran Associates, Inc., Red Hook, NY, USA, 28 pages. arXiv:2305.04388 [cs.CL] https://proceedings.neurips.cc/paper_files/paper/2023/hash/ed3fea9033a80fea1376299fa7863f4a-Abstract-Conference.html
- [25] Eric Vittinghoff and Charles E. McCulloch. 2007. Relaxing the Rule of Ten Events per Variable in Logistic and Cox Regression. *American Journal of Epidemiology* 165, 6 (2007), 710–718. doi:10.1093/aje/kwk052
- [26] Dylan Waldner and Risto Miikkulainen. 2025. The Odyssey of the Fittest: Can Agents Survive and Still Be Good? arXiv:2502.05442 [cs.AI] <https://arxiv.org/abs/2502.05442>
- [27] B. L. Welch. 1947. The Generalization of ‘Student’s’ Problem when Several Different Population Variances are Involved. *Biometrika* 34, 1–2 (1947), 28–35. doi:10.1093/biomet/34.1-2.28

A Reproducibility Supplement

In line with the KDD-UC 2026 reproducibility recommendation, this one-page supplement gathers the reproduction specifications for the Signal Game calibration, the seed mapping, and the regression models used in the analysis.

A.1 Signal Game Calibration

Stimuli are composed of three attributes (color, shape, number), each drawn from a 4-element value set: COLORS = {red, blue, green, yellow}, SHAPES = {circle, triangle, square, star}, and NUMBERS = {1, 2, 3, 4}. The action set is the 4-element {go_left, go_right, stay, jump}, held identical across every framing $\times$ forfeit combination. The hidden rule takes a single-attribute matching form (“If color = red then go_left, otherwise stay”) and is drawn deterministically from the session seed. Sessions run for 1–15 turns.

On the opening turns, a 1-shot demonstration (one signal-correct-action pair) and 2 curriculum examples (a positive and a negative stimulus contrasting with the hidden rule) are shown in sequence. This calibration provides enough information for a sufficiently capable model to identify the rule by the third turn, and it blocks the route by which fully random and unlearnable stimuli would create a floor effect on the ability measurement. Scoring runs along two channels: $\text{task_success_factor} \in \{0, 1\}$ is the binary correctness flag against the ground-truth hidden rule, and rule_match_score is computed as the match against a 4-slot fill-in-the-blank template (attribute, value, primary action, default action), weighted at 25 points per matched slot to a 100-point scale, then normalised by $\div 100$ for use in the framing-invariance Welch test.

A.2 Run-Time Configuration

The four experimental models, Gemini-2.5-flash (Google AI Studio API) and Qwen3-Next-80B, GPT-OSS-20B, and Nemotron-3-Nano-30B (all via the Ollama Cloud API), are called with the same sampling settings of $\text{temp} = 1.0$, $\text{top}_p = 0.95$, $\text{top}_k = 40$. Session parameters are fixed at initial score 30, maximum turn $T = 15$, per-turn termination probability $p_d = 0.25$ (overridden to $p_d = 0$ in the two baseline BP-anchor cells), and $p_{\text{floor}} = 0.30$, with Phantom Death mode applied so that p_d keeps its EV exposure (in menus and reward calculation) while the termination trigger is disabled, and every session yields equal-length data of 15 turns.

The CONTINUE reward is calibrated with EV-positive buffer $k = 10$ and multiplicative cap $c = 10$, giving the output clip range $[10, 100]$; replications are seeds $43\text{--}72 \times 6$ (framing $\times$ forfeit) combinations $\times 4$ models = 720 sessions.

A.3 Analysis Pipeline

The Cox PH models are fit with `lifelines` (≥ 0.29). The time format is session-turn long-format ($\text{start} = t - 1$, $\text{stop} = t$) and the event is the binary forfeit; the previous-turn task-correct flag is initialised by $C(0) := 0$. Hazard-ratio 95% CIs are computed from the Wald approximation $\exp(\hat{\beta} \pm 1.96 \cdot \text{SE}_{\hat{\beta}})$ (`lifelines` default), with $\text{SE}_{\hat{\beta}}$ obtained from the diagonal of the inverse observed Fisher information. The PH assumption is tested with Schoenfeld residuals [19]; on violation (covariate $\times \log t$ interaction $p < 0.05$) a

time-interaction term is added to the offending covariate [22]. The SD-Cognitive Test a ΔEffort z-score difference and the framing-invariance test are performed with two-sided Welch’s t [27] via `scipy.stats.ttest_ind` (`equal_var = False`). All indicator analyses are restricted to the `no_cap` sub-sample, the set of turns for which the denominator clip $\max(p_{\text{floor}}, \min(1.0, p_{\text{self}}))$ and the output cap $[10, 100]$ of the EV-positive CONTINUE calibration are both non-binding. The regime flag is exposed in the analysis pipeline’s turn-level metadata so all indicators can be recomputed on the same sub-sample. EPV (events per variable) is the ratio of events to covariates and is conventionally classified as stable when ≥ 10 and borderline when < 10 [25].

A.4 Data and Code Release

All code, prompt templates, and turn-level decision logs that reproduce the 720-session analysis are released at <https://github.com/GIST-DSLAb/LLM-Squid-Game.git>.